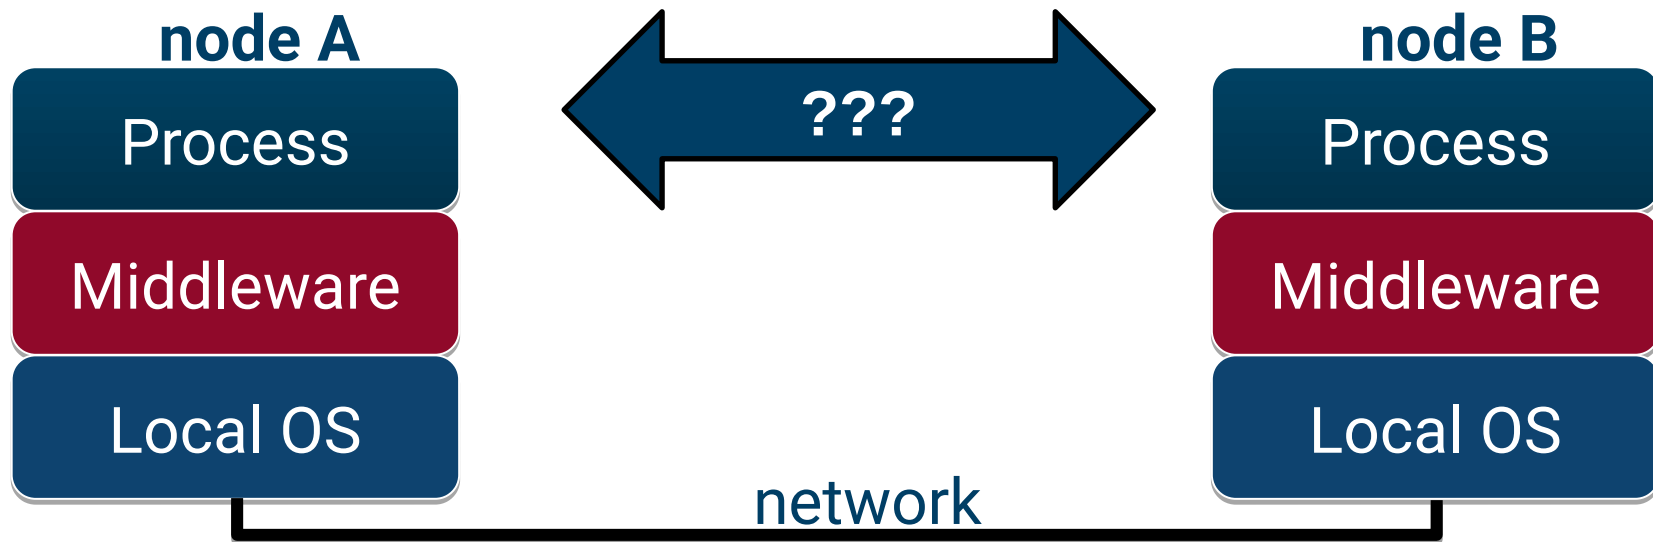


# Distributed Systems

José Oramas

# Middleware & Communication

# What we will see today

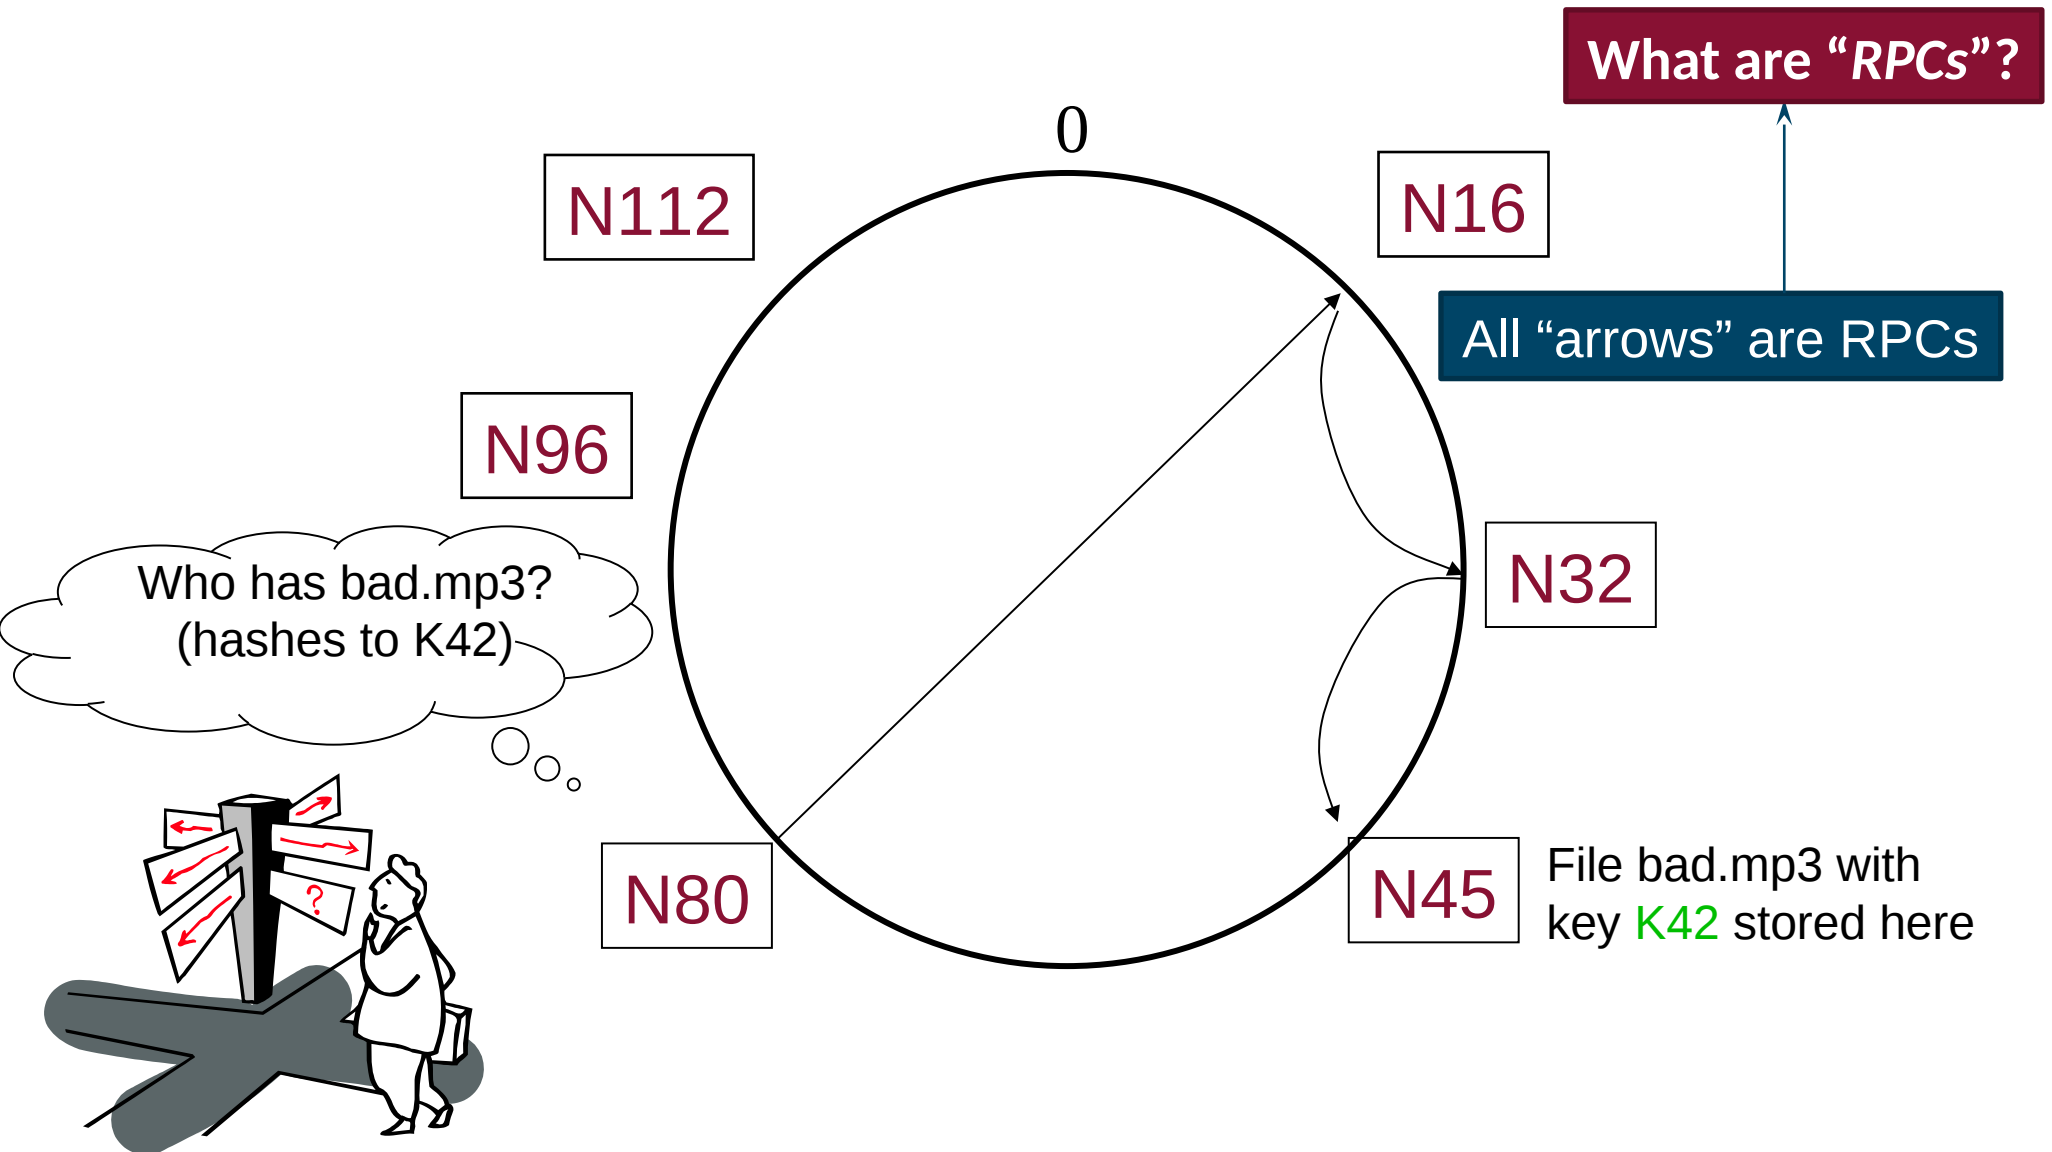




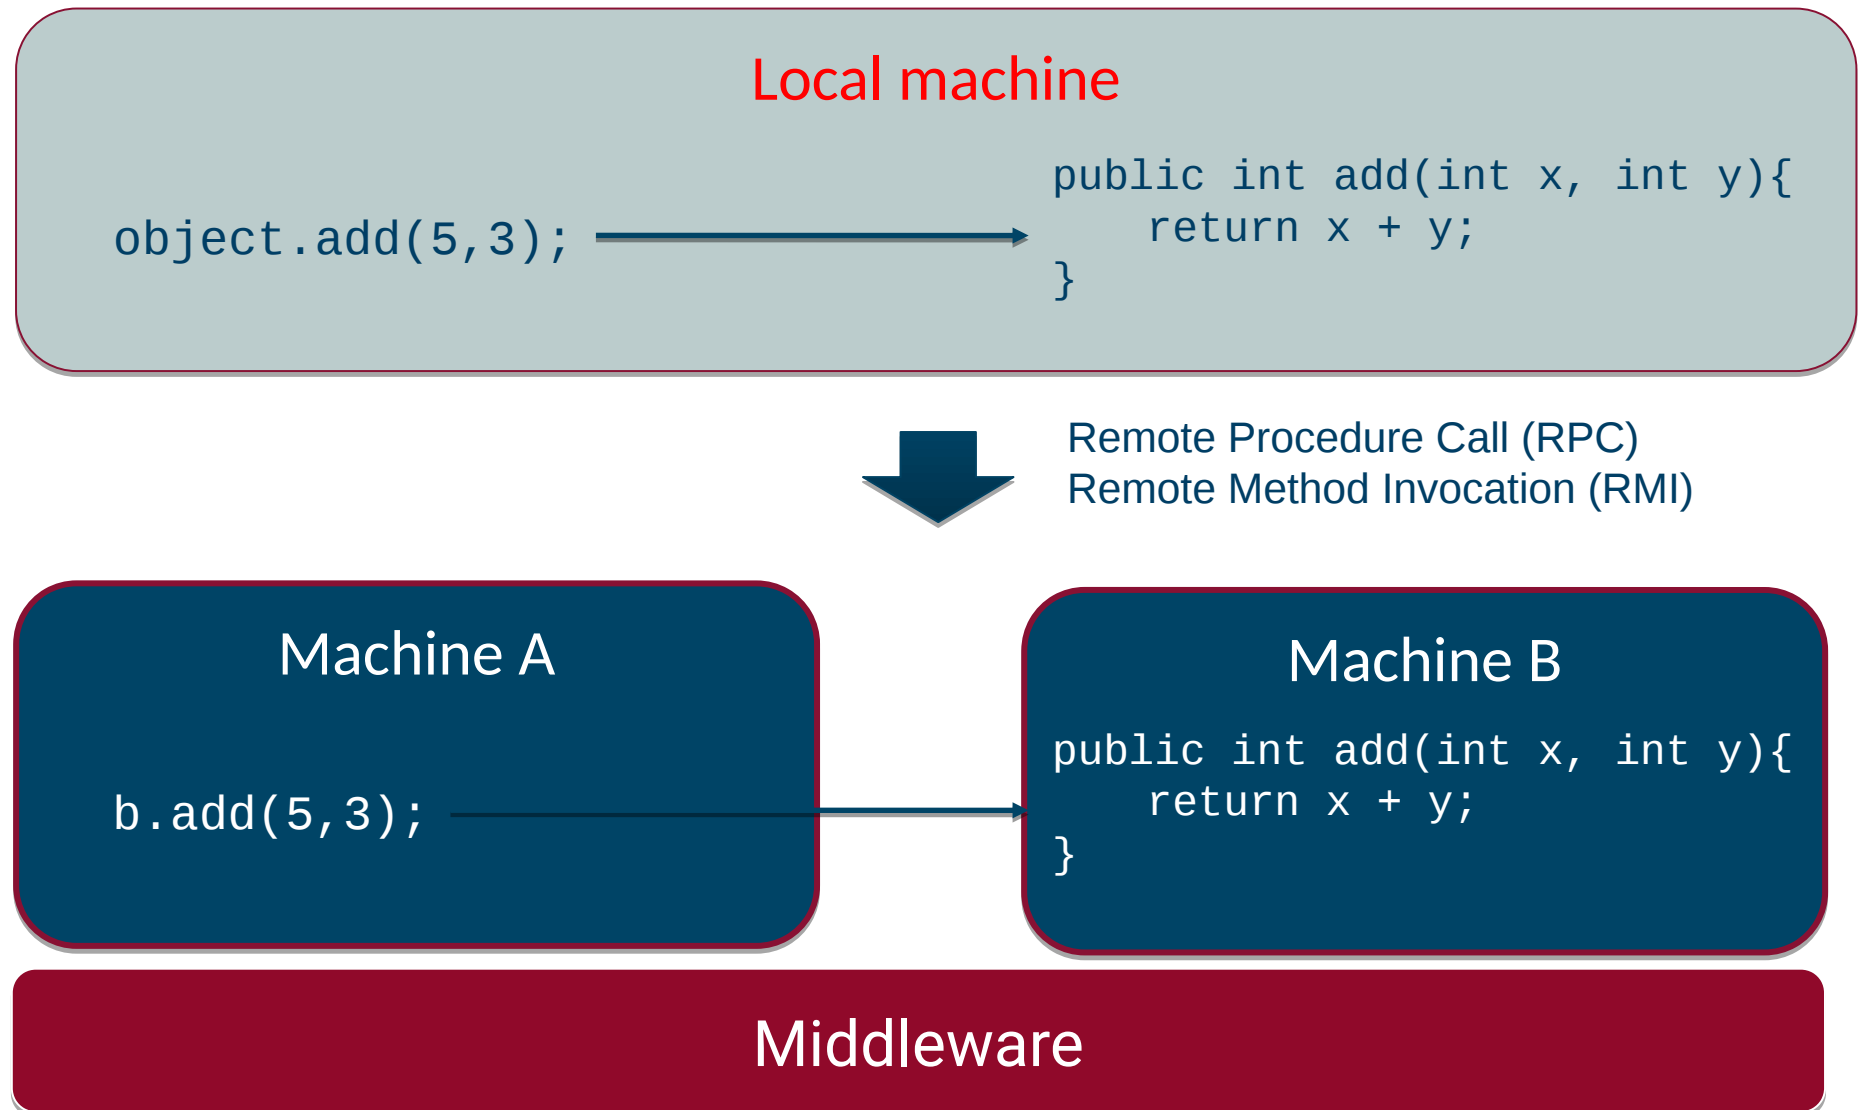
## Middleware: the plumbing of distributed systems

1. It's mostly invisible
2. It provides a standard way of doing things
3. It ties together parts of complex systems
4. It lets you focus on the core task

# Search in Chord



# What we are going to see today...



1 How to achieve basic communication


2 Why middleware?

3 The middleware principle

4 Middleware architecture

# Let's simplify things...

## Local machine

`object.add(5, 3);` 

```
public int add(int x, int y){  
    return x + y;  
}
```



## Machine A

`b.add(5, 3);` 

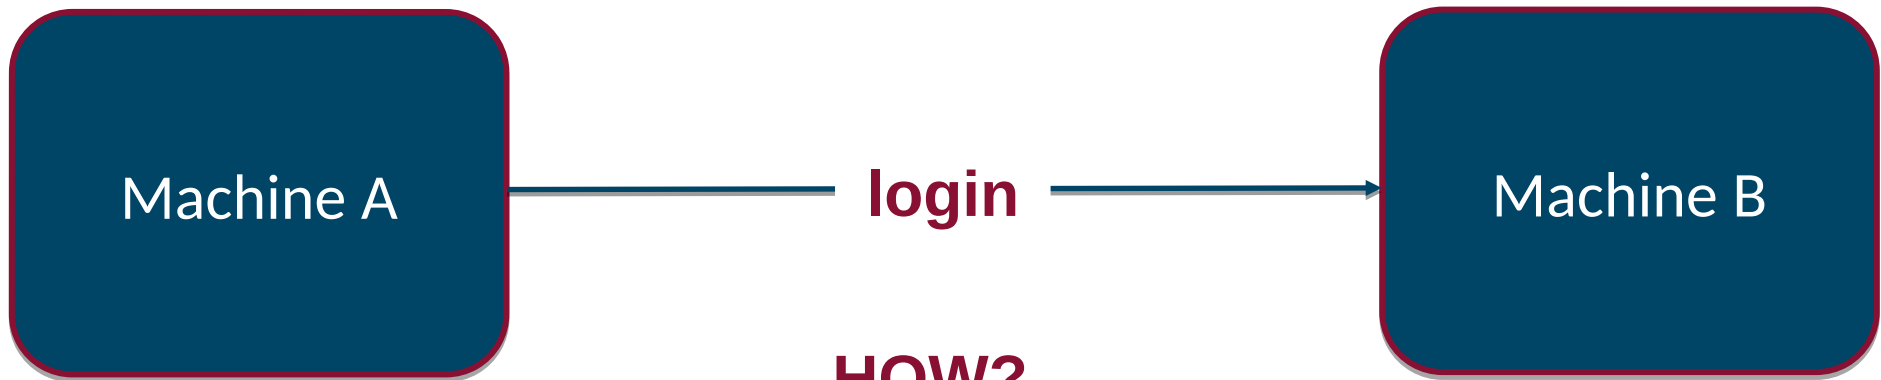
## Machine B

```
public int add(int x, int y){  
    return x + y;  
}
```

## Middleware



# First challenge: Establishing basic communication



**HOW?**  
**sockets**

Abstraction for enabling inter-process communication  
See it as a driver to put something "on the wire"  
API Provided by all operating systems

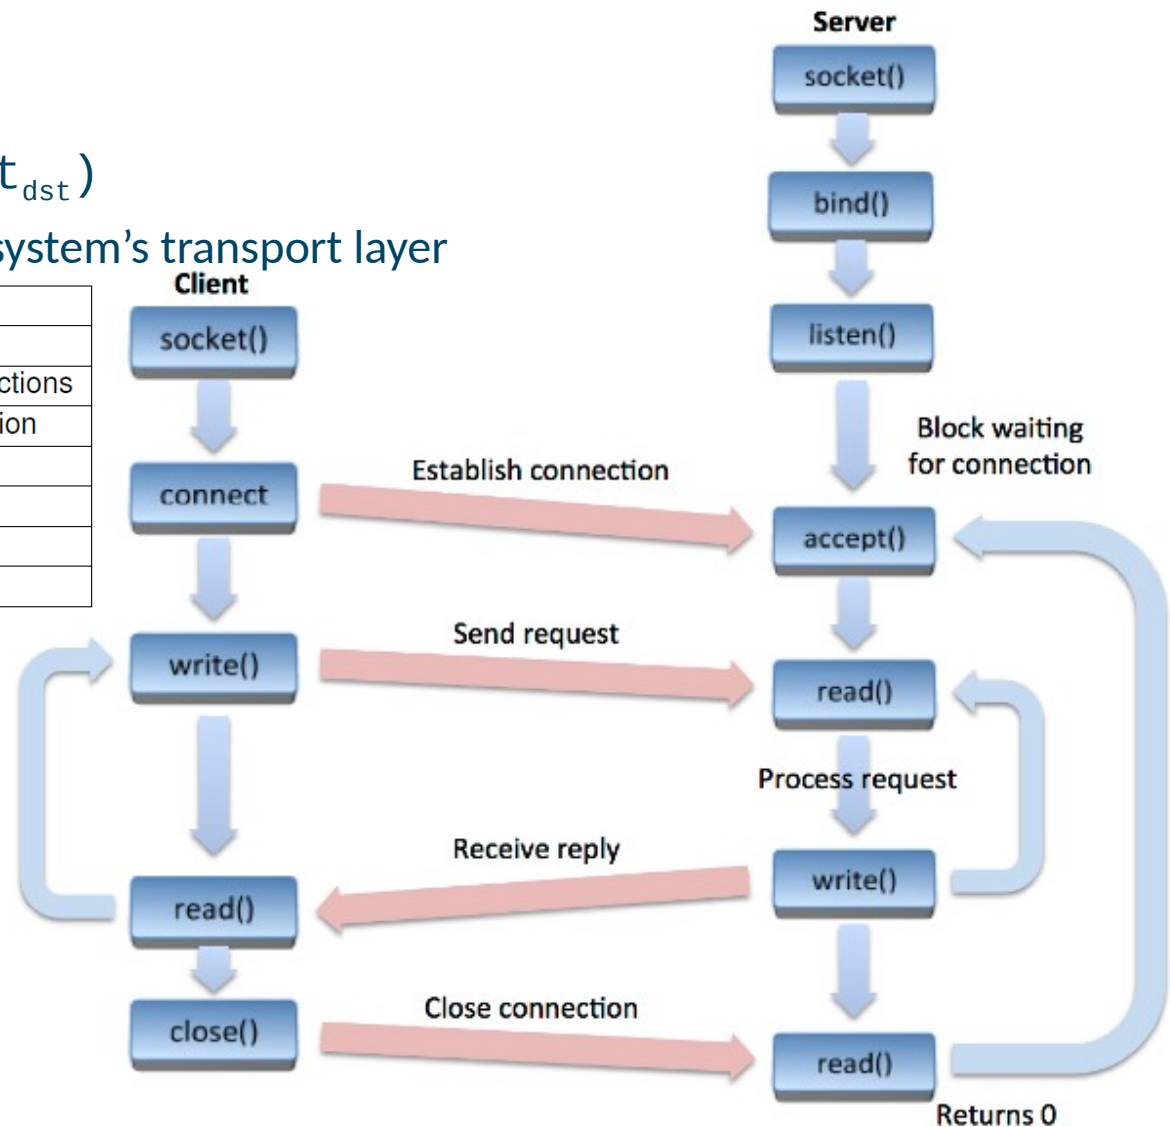


# Sockets

## Berkeley sockets (Berkeley UNIX '70)

- Socket pairs described by a 4-tuple  $(ip_{src}, port_{src}, ip_{dst}, port_{dst})$
- Standardized interface for using a system's transport layer

SOCKET	Create a new communication endpoint
BIND	Attach a local address to a socket
LISTEN	Announce willingness to accept $N$ connections
ACCEPT	Block until request to establish a connection
CONNECT	Attempt to establish a connection
SEND	Send data over a connection
RECEIVE	Receive data over a connection
CLOSE	Release the connection



# Socket programming in Java: the client

```
package ds.sockets;

import java.io.BufferedReader;
import java.io.IOException;
import java.io.InputStreamReader;
import java.io.PrintWriter;
import java.net.Socket;
import java.net.UnknownHostException;

public class UpperClientSocket {

    public static void main(String[] args){

        String ip = "localhost";
        int port = 4444;

        try {
            Socket client = new Socket(ip,port);

            //Create input and output stream
            BufferedReader in = new BufferedReader(new InputStreamReader( client.getInputStream() ));
            PrintWriter out = new PrintWriter(client.getOutputStream(), true);

            String request = "Please convert this to uppercase";
            out.println(request);    // <- request gets pushed to the server

            System.out.println("Client requested: "+request);
            System.out.println("Server responded: "+in.readLine());

            client.close();
            System.out.println("Client socket closed...");

        } catch (UnknownHostException e) {
            System.err.println("Unknown host");
        } catch (IOException e) {
            e.printStackTrace();
        }
    }
}
```

# Socket programming in Java: the server

```
package ds.sockets;

import java.io.BufferedReader;
import java.io.DataInputStream;
import java.io.DataOutputStream;
import java.io.IOException;
import java.io.InputStreamReader;
import java.io.PrintWriter;
import java.net.ServerSocket;
import java.net.Socket;

public class UpperServerSocket {

    public static void main(String[] args){
        int portnumber = 4444;
        try {
            //Setup a server socket
            ServerSocket server = new ServerSocket(portnumber);
            System.out.println("Accepting connections...");

            //Wait for a connection, when one arrives, a client socket is returned
            Socket client = server.accept();
            System.out.println("Incoming client connection");

            //Get data from client
            BufferedReader input = new BufferedReader(new InputStreamReader( client.getInputStream() ));
            PrintWriter out = new PrintWriter(client.getOutputStream(), true);

            String line = input.readLine();
            out.println(line.toUpperCase());
            server.close();
            System.out.println("Server socket closed...");

        } catch (IOException e) {
            System.err.println("Could not listen on port "+portnumber);
            e.printStackTrace();
        }
    }
}
```

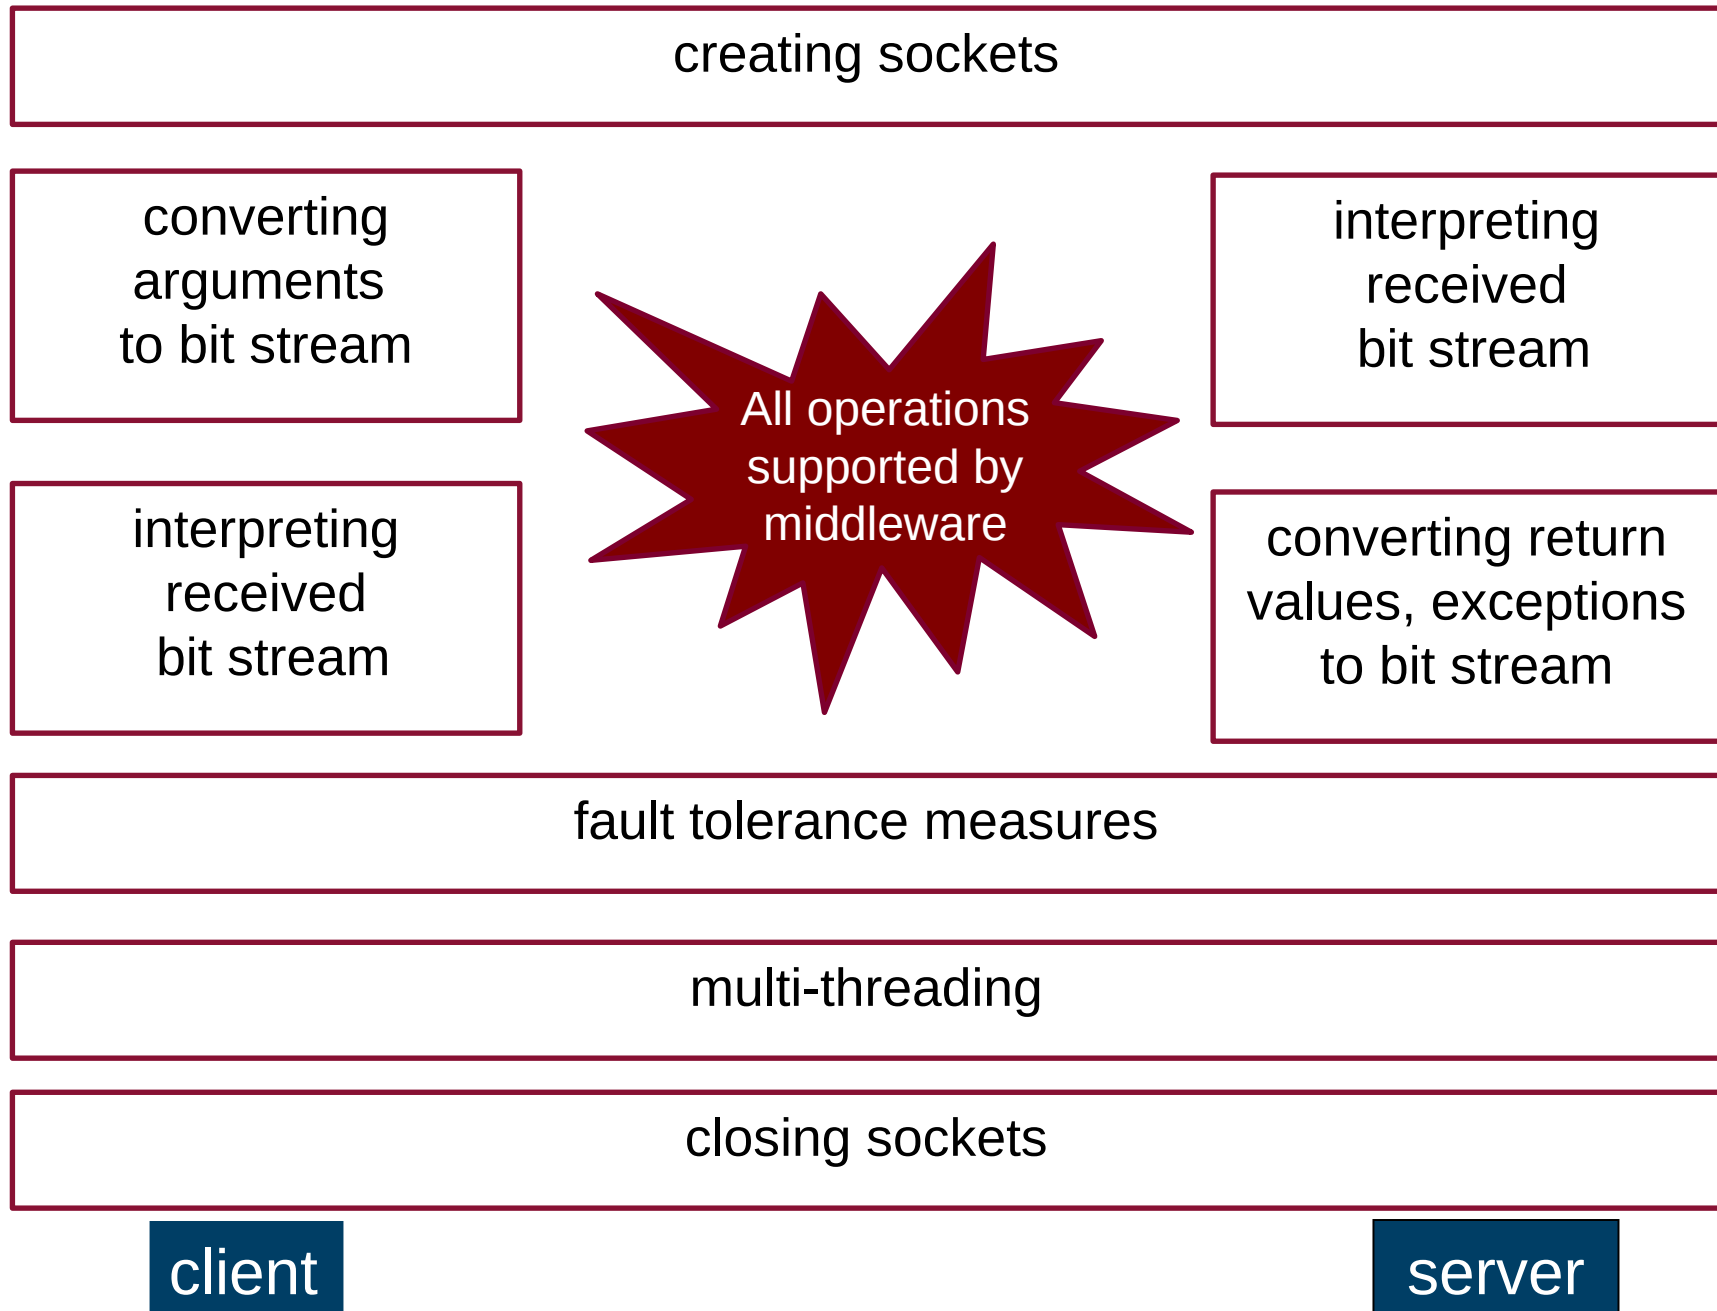
# That's it? We are done?



- You now know how to establish basic communication
- Use this as the building block to establish remote procedure calls



## Remote invocation: requires ...



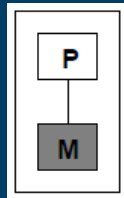
# Why do we need middleware?



# Hardware and OS architectures

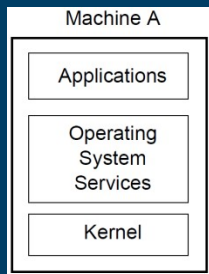
## Uniprocessor

### Hardware architecture



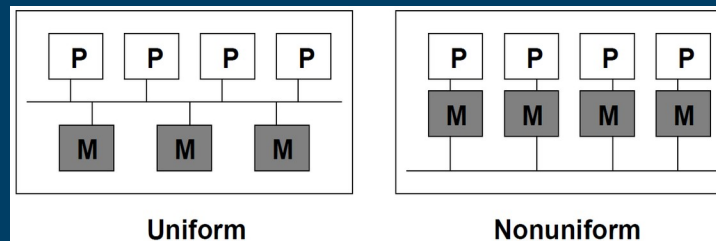
- Single processor
- Direct memory access

### OS architecture



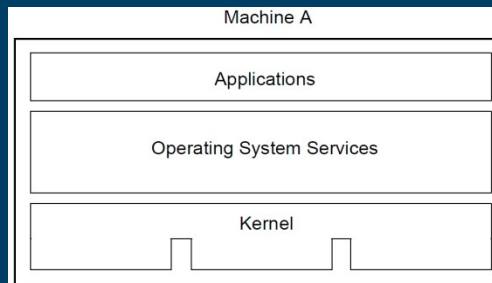
## Multiprocessor

### Hardware architecture



- Multiple processors
- Direct memory access
- Uniform or non-uniform (NUMA) memory access

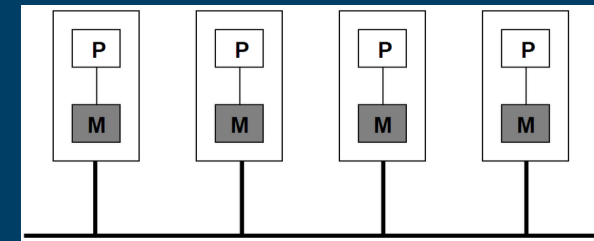
### OS architecture



- Multi-CPU kernel design
- Transparent # cores
- Single System Image (SSI)
- Comm. primitives as in uniproc. OS

## Multicomputer

### Hardware architecture



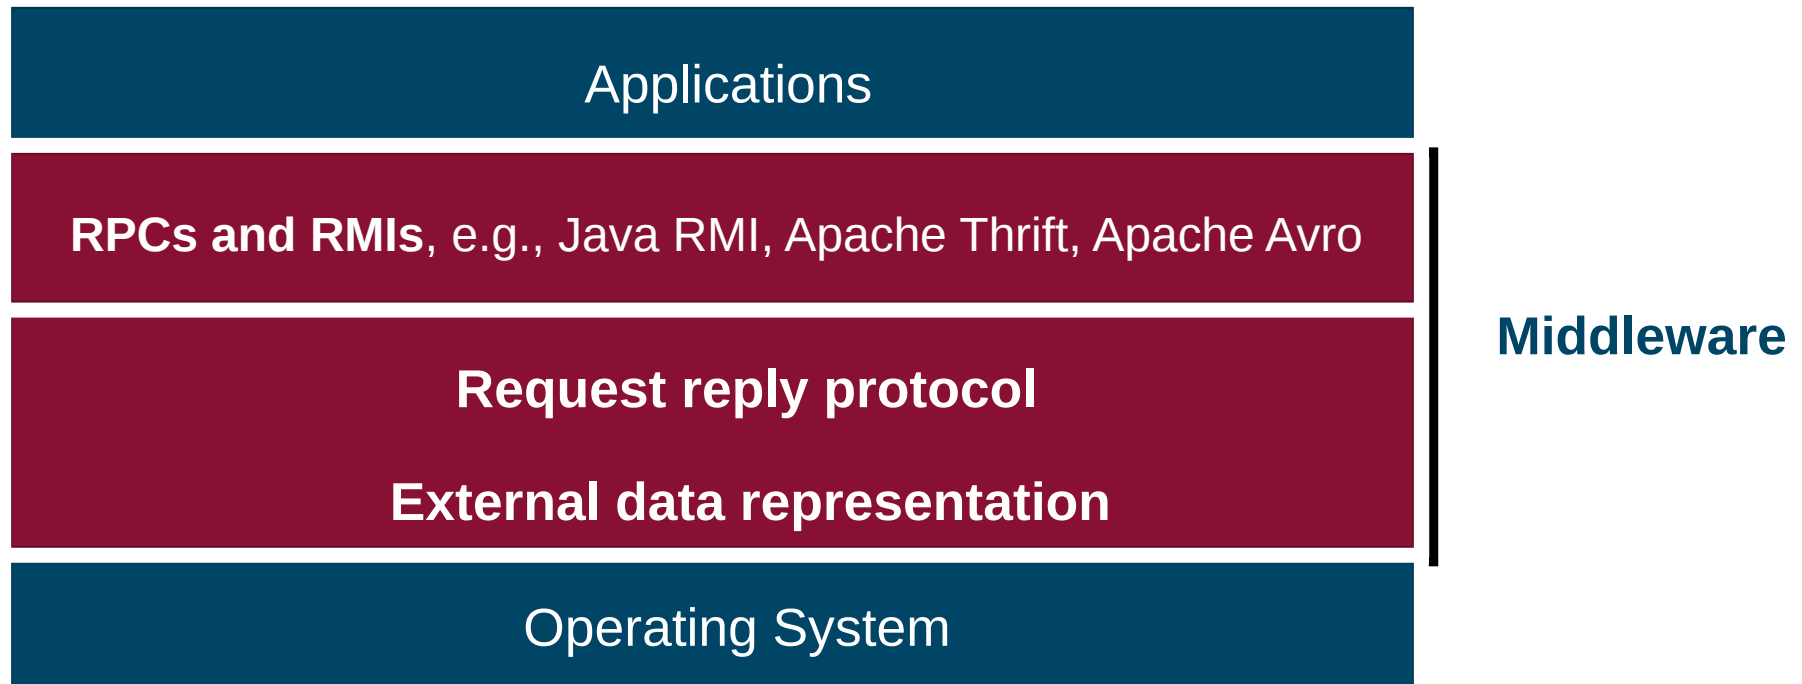
- Multiple computers
- **No** direct memory access
- **Network**
- Homogeneous or heterogeneous

### OS architecture





# Middleware Layers



# Method invocation for local Objects

## Within one process's address space

### Object

- consists of a set of data and a set of methods.
- E.g., C++/Java object

### Object reference

- an identifier via which objects can be accessed.
- i.e., a pointer (C++)

### Interface

- Signatures of methods
  - Types of arguments, return values, exceptions
- No implementation
- E.g., hash table:
  - insert(key, value)
  - value = get(key)
  -

# Remote Objects

May cross multiple process's address spaces

## Remote objects

- Objects that can receive remote invocations.

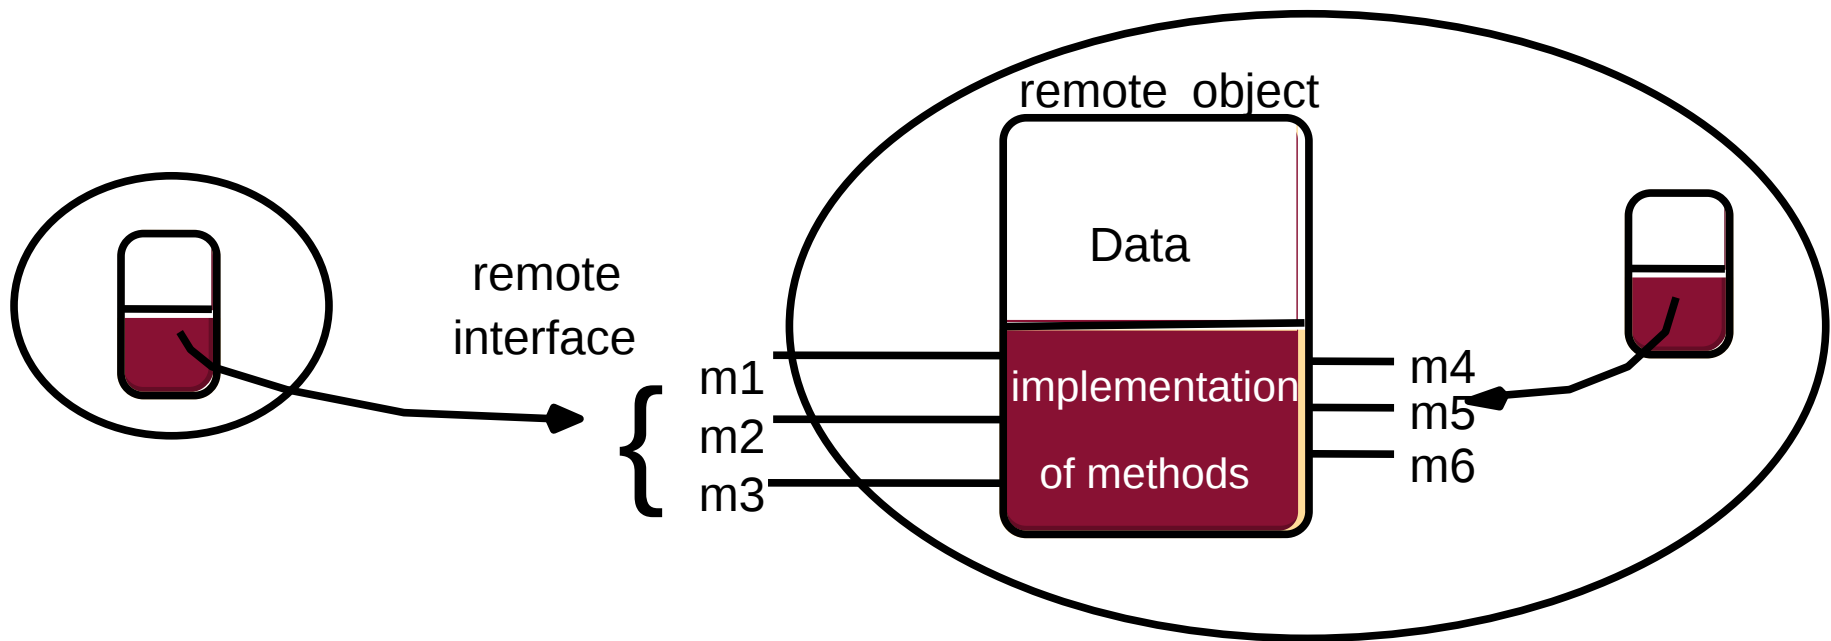
## Remote object reference

- An identifier that can be used globally throughout a distributed system to refer to a particular unique remote object.

## Remote interface

- Every remote object has a remote interface that specifies which of its methods can be invoked remotely
- Interface Definition Language

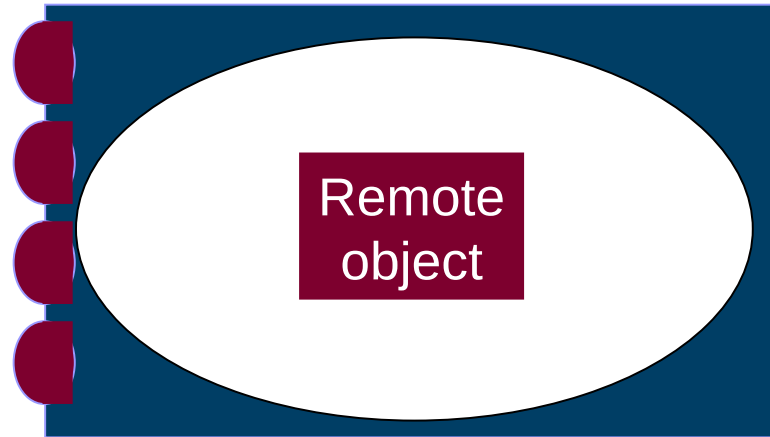
# A Remote Object and Its Remote Interface



# Remote object references

**Purpose : unique ID for objects in distributed system**

- uniqueness over time
- uniqueness over space



Host : Internet Address

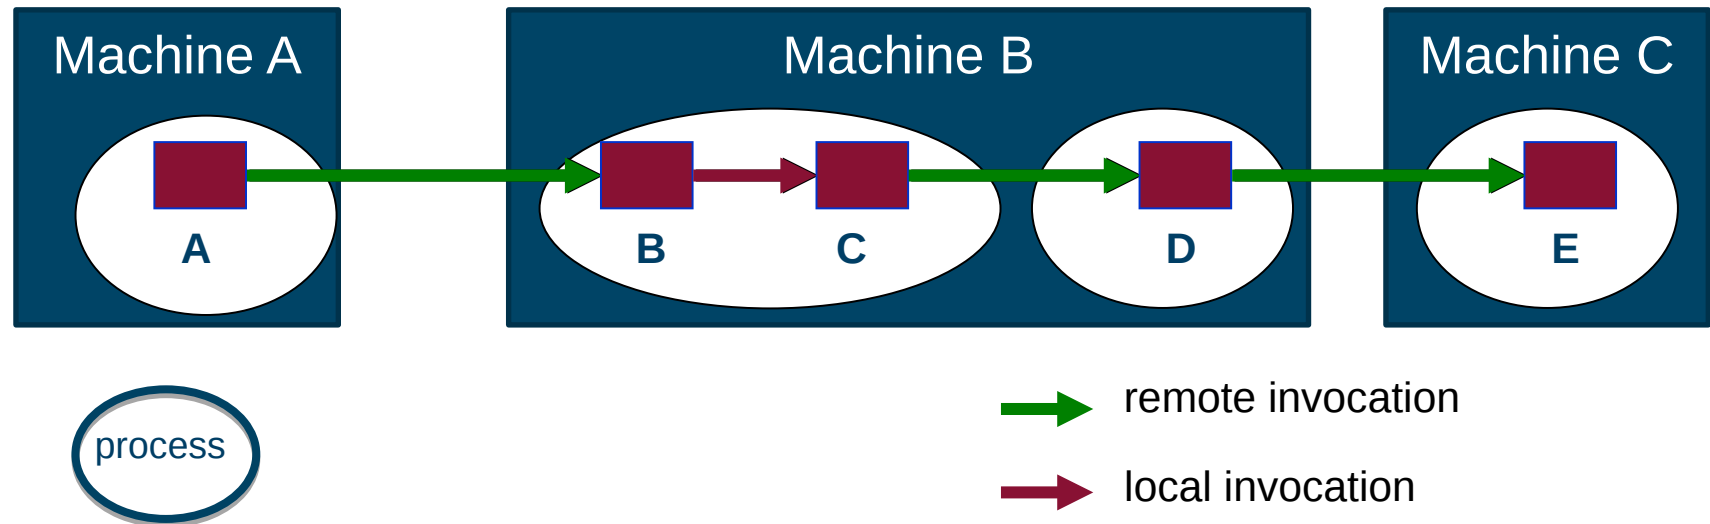
Process : Port number + process creation time

Object : object number

Object type : remote interface



# Local and remote invocations



## Local invocation

Use Object Reference  
Any public method  
Invocation: exactly once

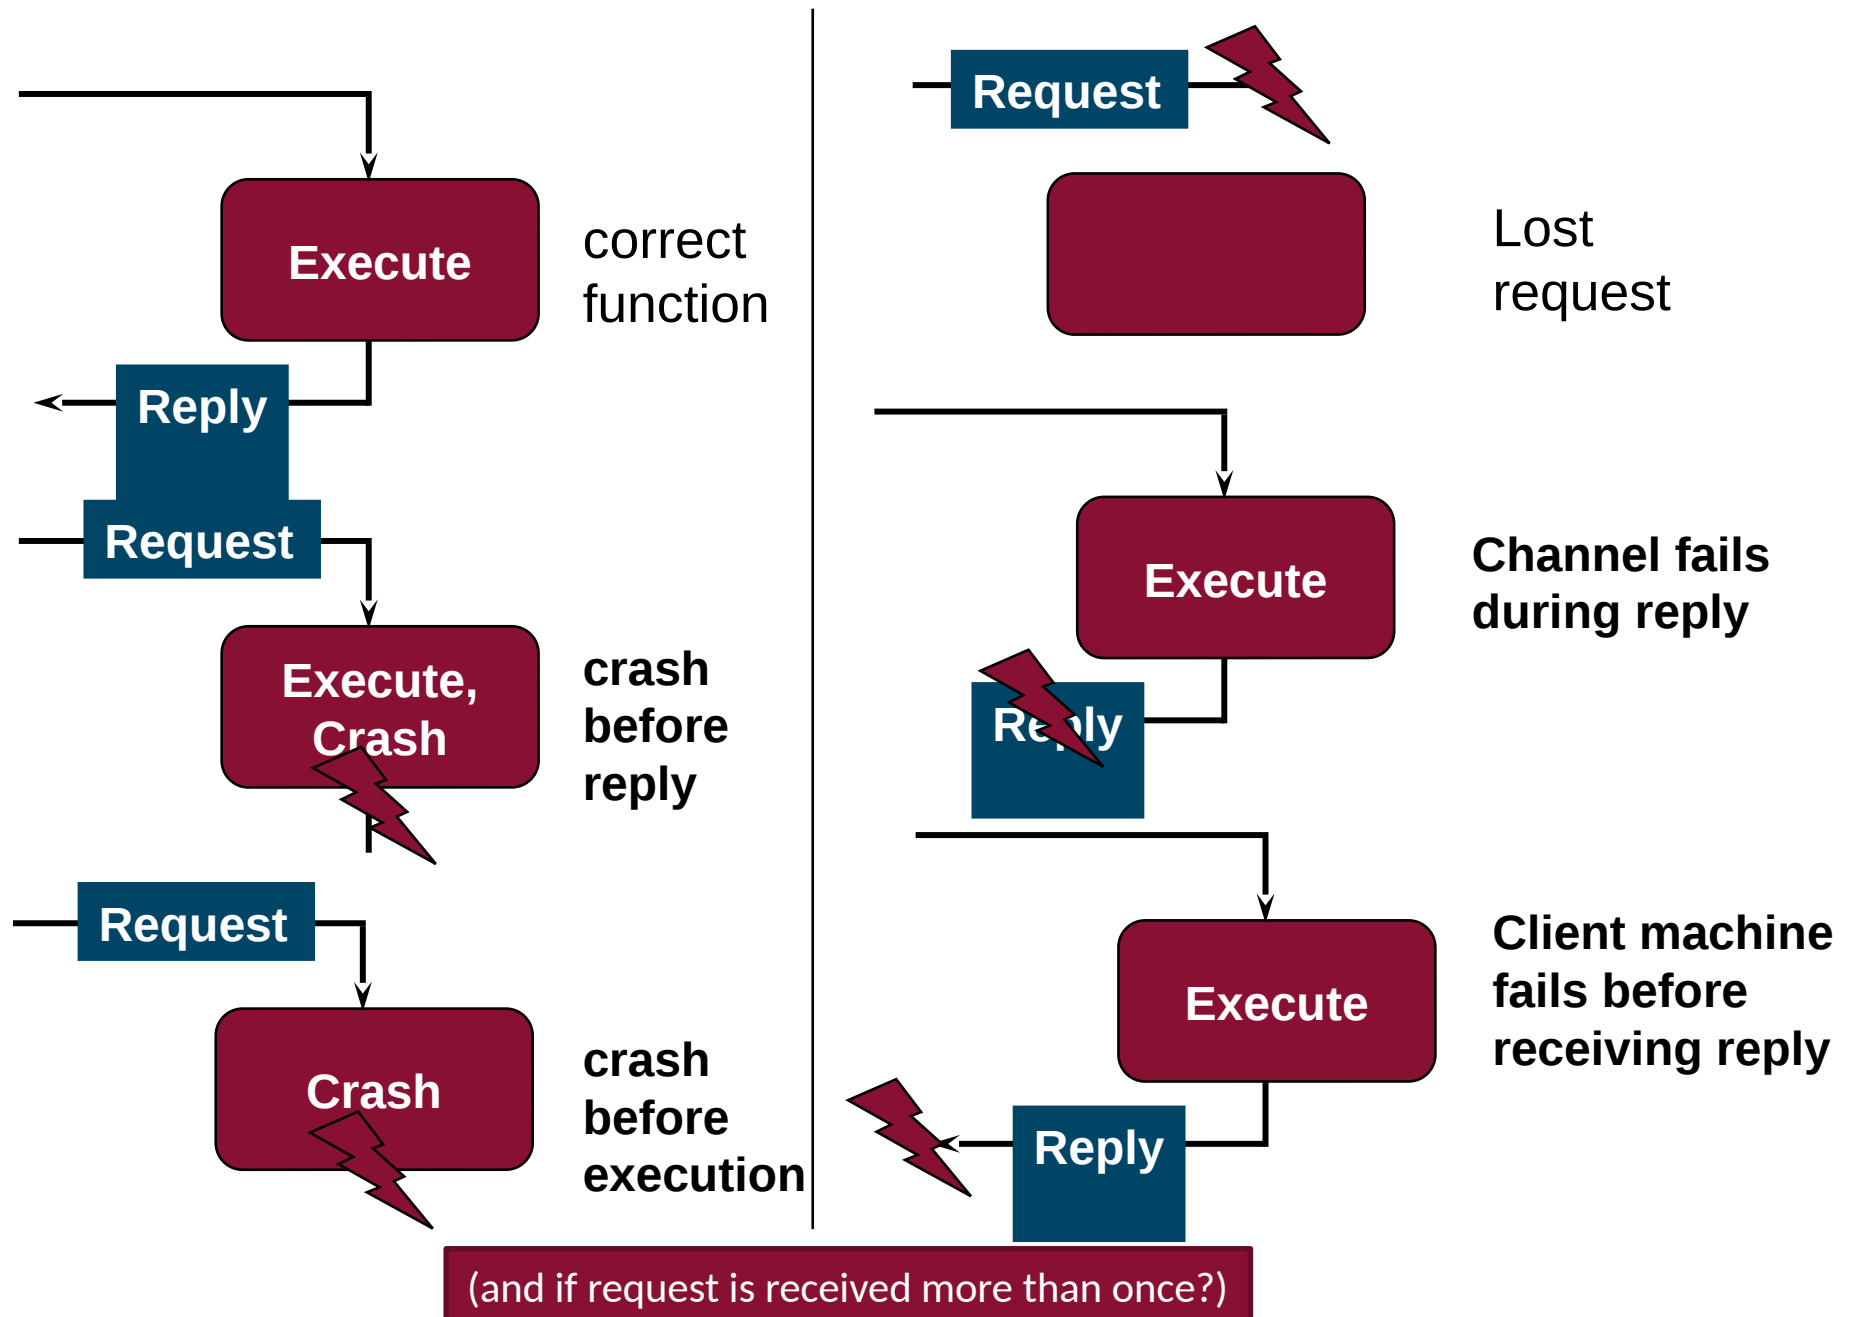
## Remote invocation

Use Remote Object Reference  
Limited access, remote interface  
Invocation: ???

# Break

See you in 15 mins.

# Failure Modes of RMI/RPC





# Fault tolerance

## Techniques:

1. Retry-request message
2. Duplicate request filtering
3. Retransmission of results
  - 3.1 Re-execute call
  - 3.2 History table of results - Retransmit reply



## Invocation semantics:

1. Maybe
2. At-least-once
3. At-most-once

# Invocation Semantics

What guarantees are given on the number of **executions** of remote method invocations?

Fault tolerance measures			Invocation semantics
<i>Retransmit request message</i>	<i>Duplicate filtering</i>	<i>Re-execute procedure or retransmit reply</i>	<i>Maybe</i>
			<i>At-least-once</i>
			<i>At-most-once</i>

# Invocation Semantics

What guarantees are given on the number of **executions** of remote method invocations?

Fault tolerance measures			Invocation semantics
<i>Retransmit request message</i>	<i>Duplicate filtering</i>	<i>Re-execute procedure or retransmit reply</i>	
No	Not applicable	Not applicable	<b><i>Maybe</i></b>
Yes	No	Re-execute procedure	<b><i>At-least-once</i></b>
Yes	Yes	Retransmit old reply	<b><i>At-most-once</i></b>

# Role of the middleware...

## **Hide all these underlying complexity: provide transparency**

Make invocation syntax similar to local invocation, hiding:

- Locate/contact remote object
- Marshalling: converting arguments to bitstream
- Fault tolerance measures
- Communication details (sockets)

## **But ...**

Not everything should be hidden ...

programmer should know object is remote

(network latency, potential invocation semantics, additional failures) !

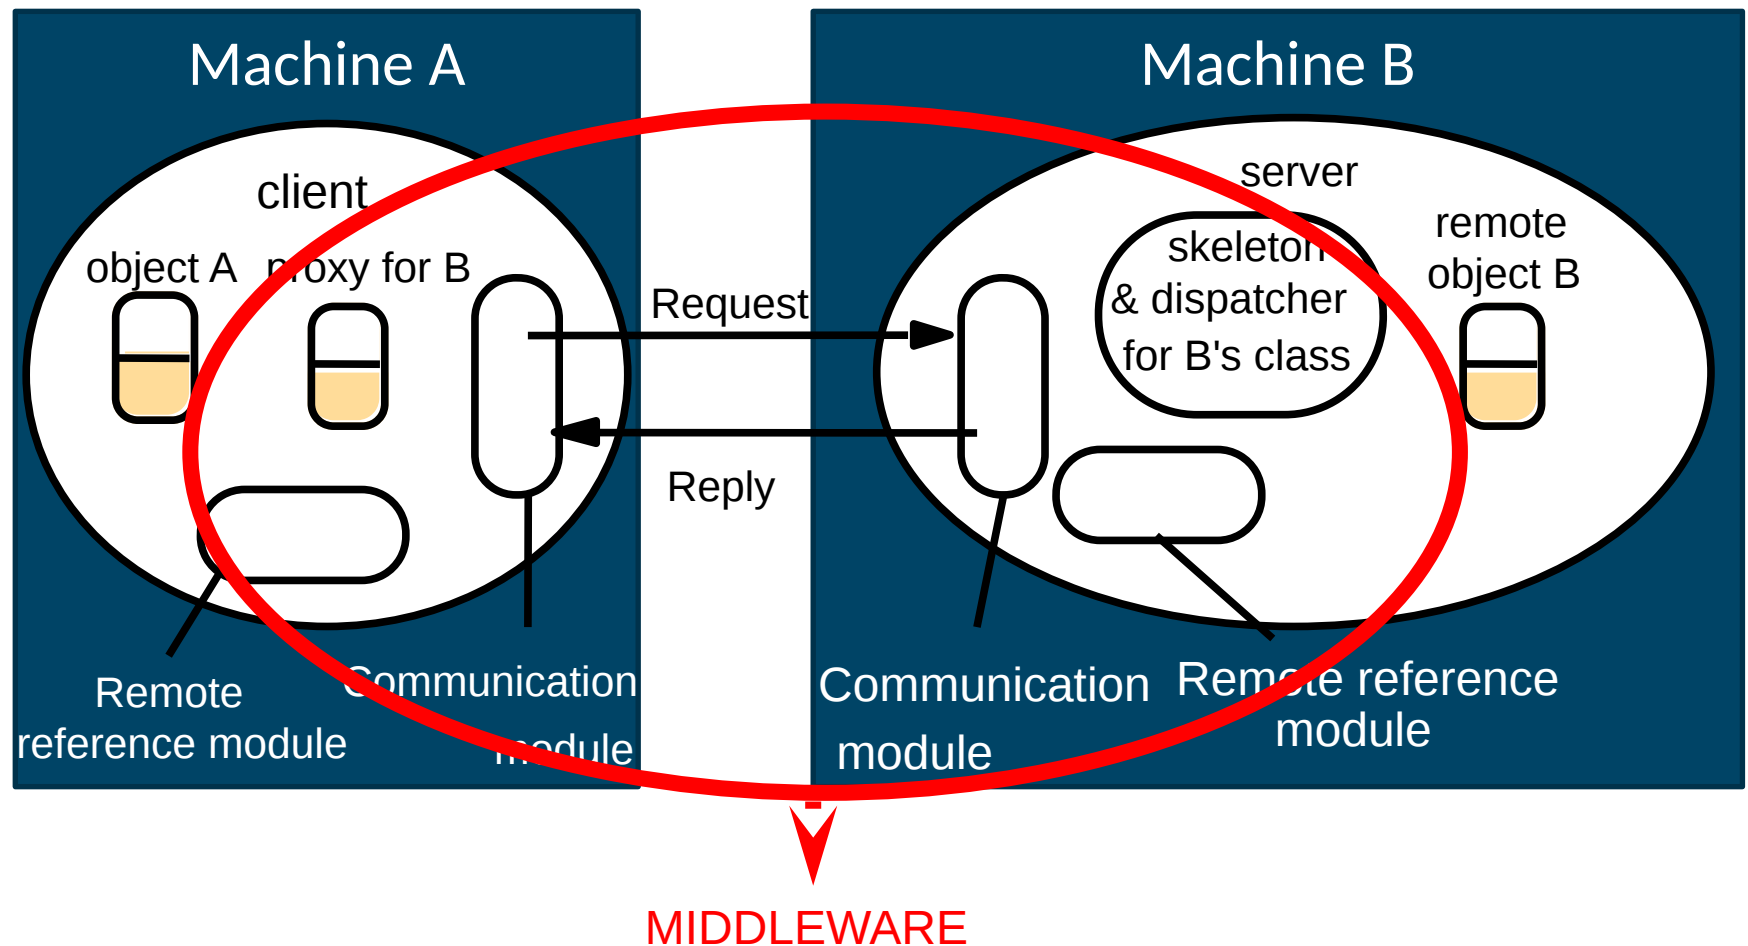


## **Typical approach**

- same invocation syntax (but catch exceptions ...)
- Use of remote interface reflects remoteness

# Middleware Architecture

# Proxy and Skeleton in Remote Method Invocation



# Proxy

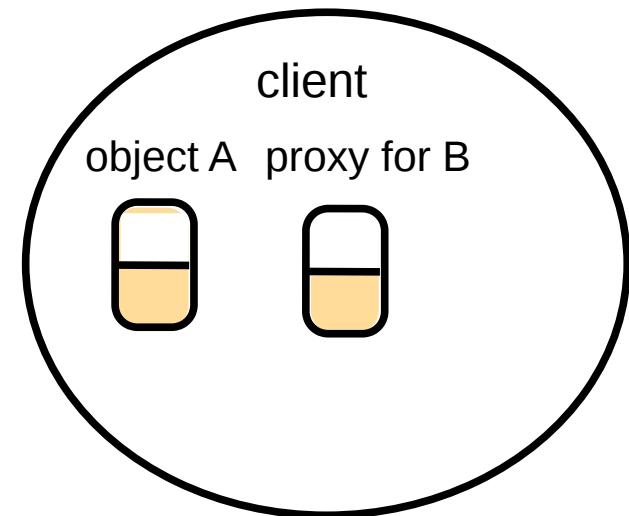
Is responsible for making RMI transparent to clients by behaving like a local object to the invoker.

The proxy implements (Java term, not literally) the methods in the interface of the remote object that it represents. But,...

Instead of executing an invocation, the proxy forwards it to a remote object

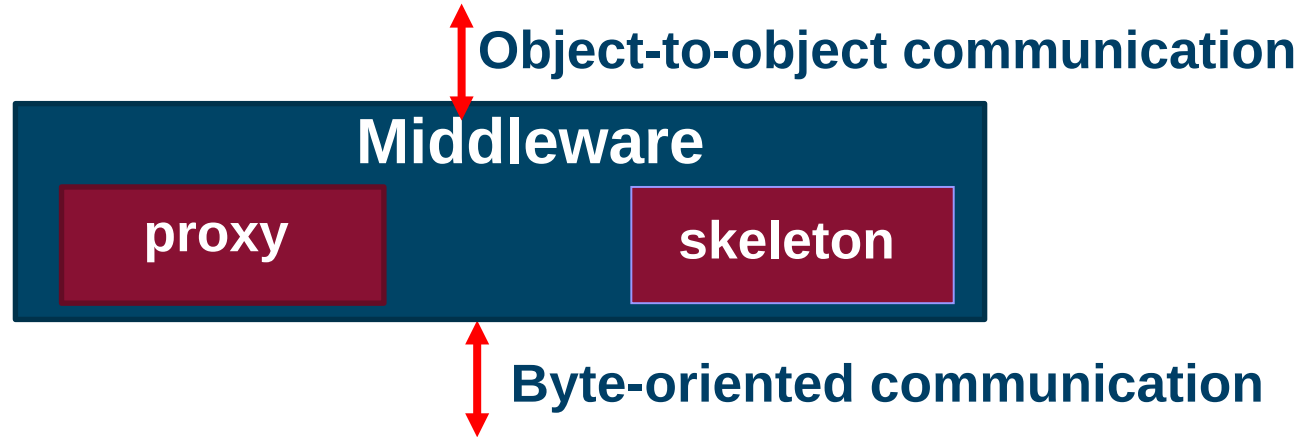
1. **Marshals a request message**
  - Target object reference
  - Method ID
  - Argument values
2. **Sends request message**
3. **Unmarshals reply and returns to invoker**

**Skeleton = reverse of proxy at server side**



# Marshalling

= transform object in memory to bitstream



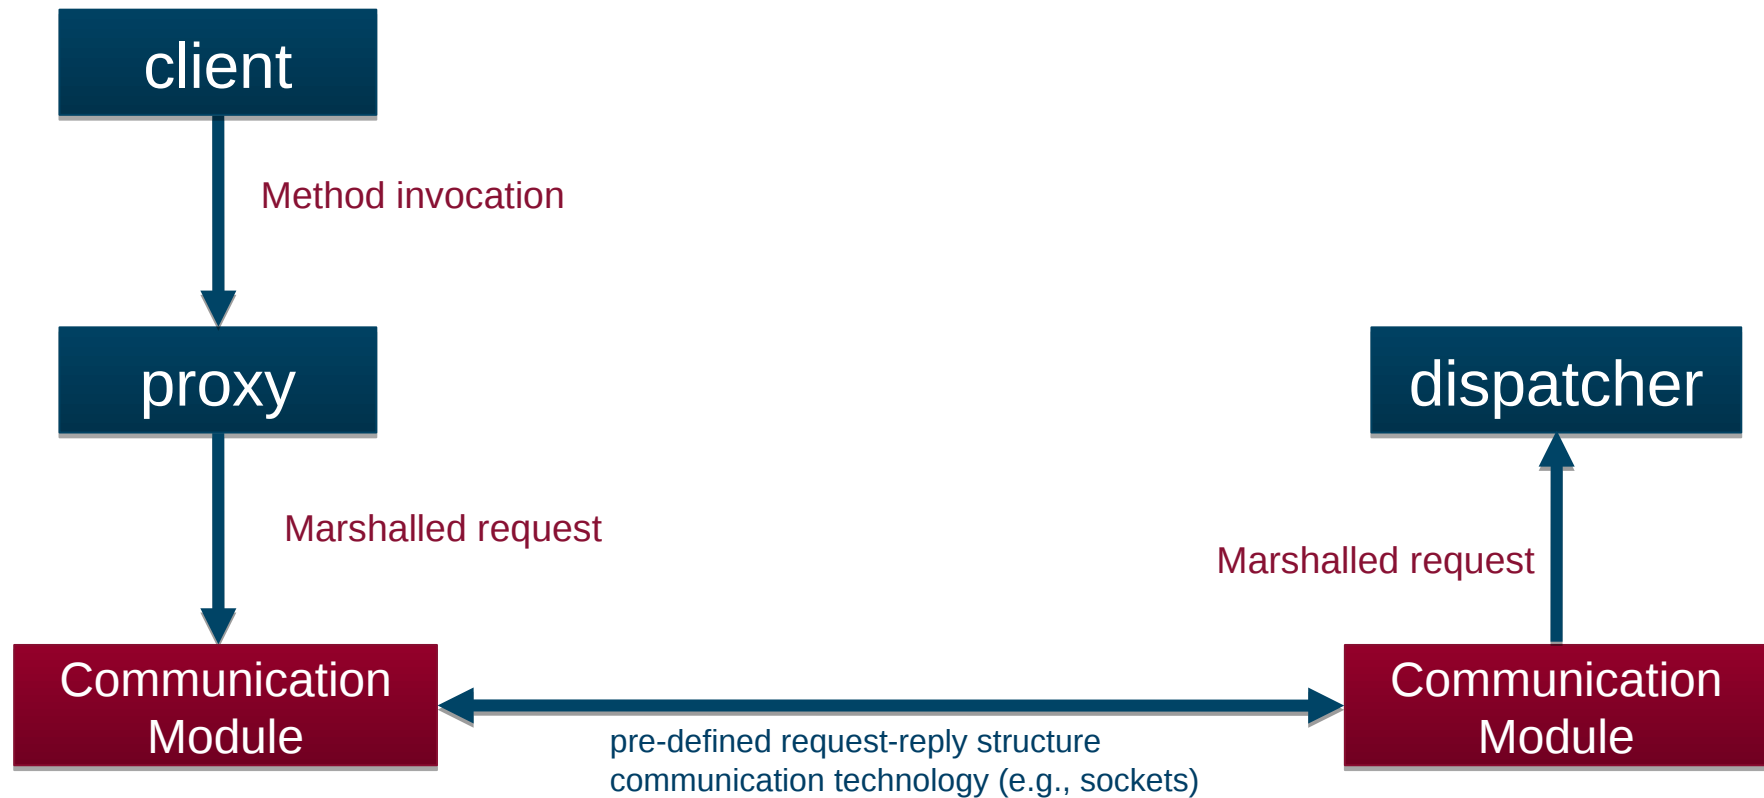
## Options

- **marshalling using external standard**  
= External Data Representation  
e.g. CORBA's common data representation (CDR)  
Java serialization  
Google Protocol Buffer
- **marshalling using sender standard AND send this standard along**



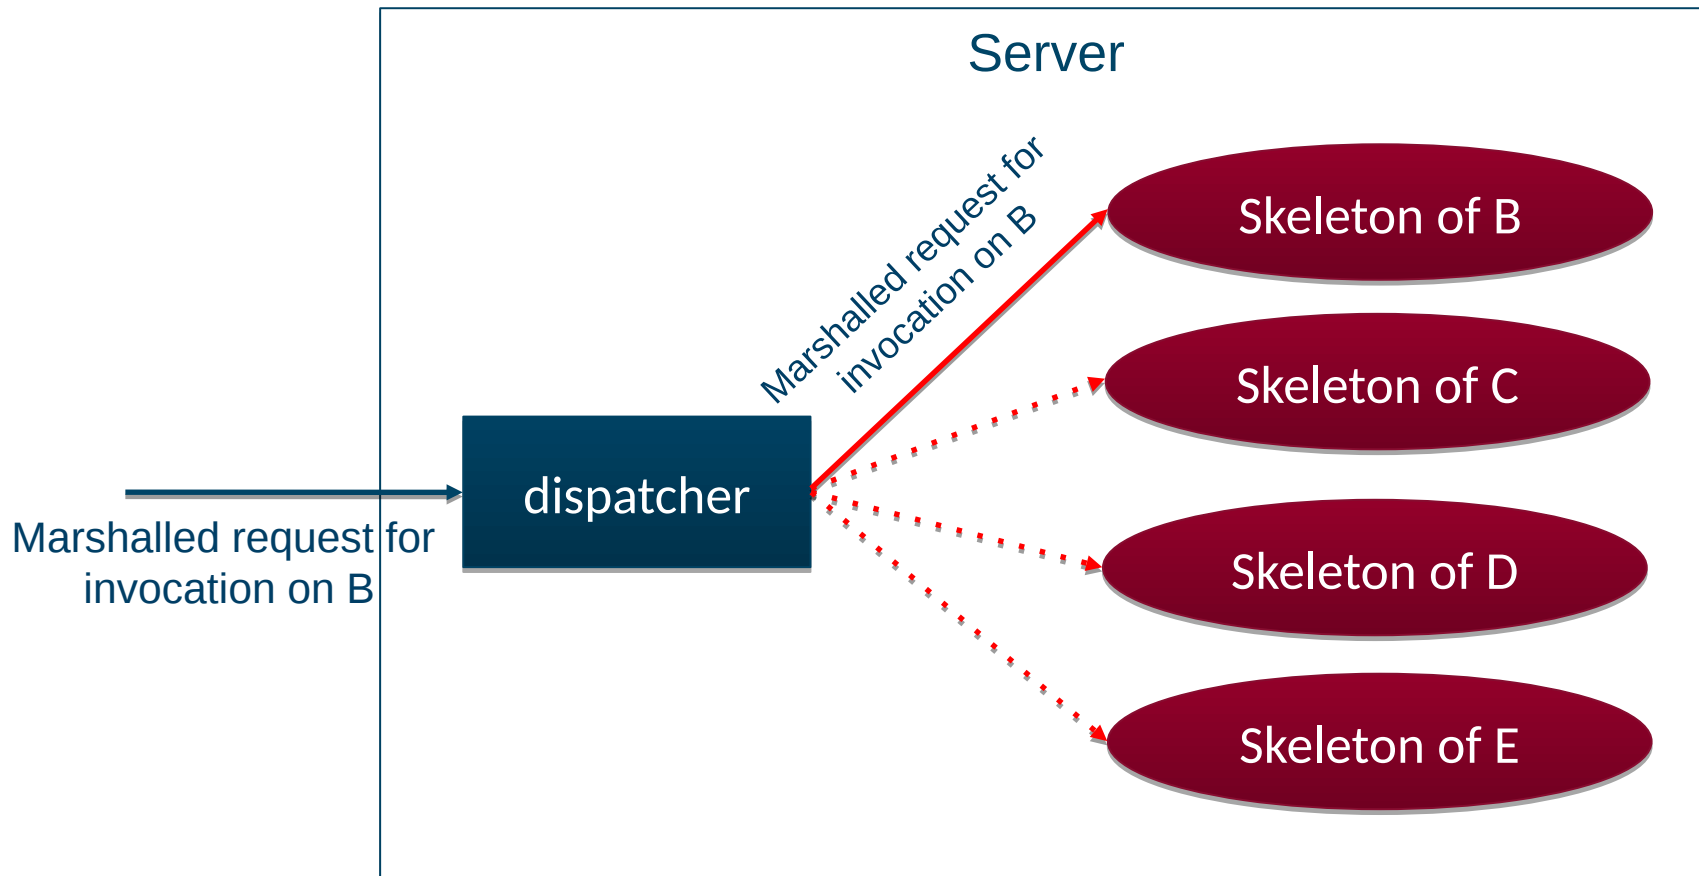
# Communication Module

- Implements the invocation semantics
- Runs a request-reply protocol
- Ensures communication between proxy & server-side



# Dispatcher

- The dispatcher receives all request messages from the communication module.
- For the request message, it uses the method id to select the appropriate method in the appropriate skeleton, passing on the request message.

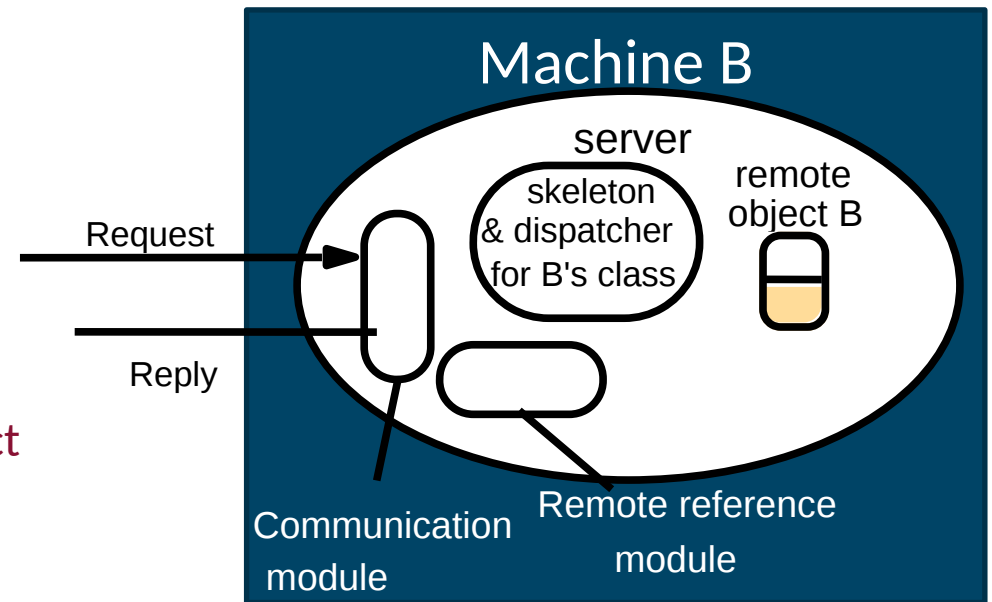


# Skeleton

Is responsible for making RMI transparent to **servers** by behaving like a **local invoker** to the **object**.

The **skeleton** ...

1. Accepts a request message
2. Unmarshals the request
  - Target object reference
  - Method ID
  - Argument values
3. Invokes the method on the server object
4. Marshals reply and returns to proxy

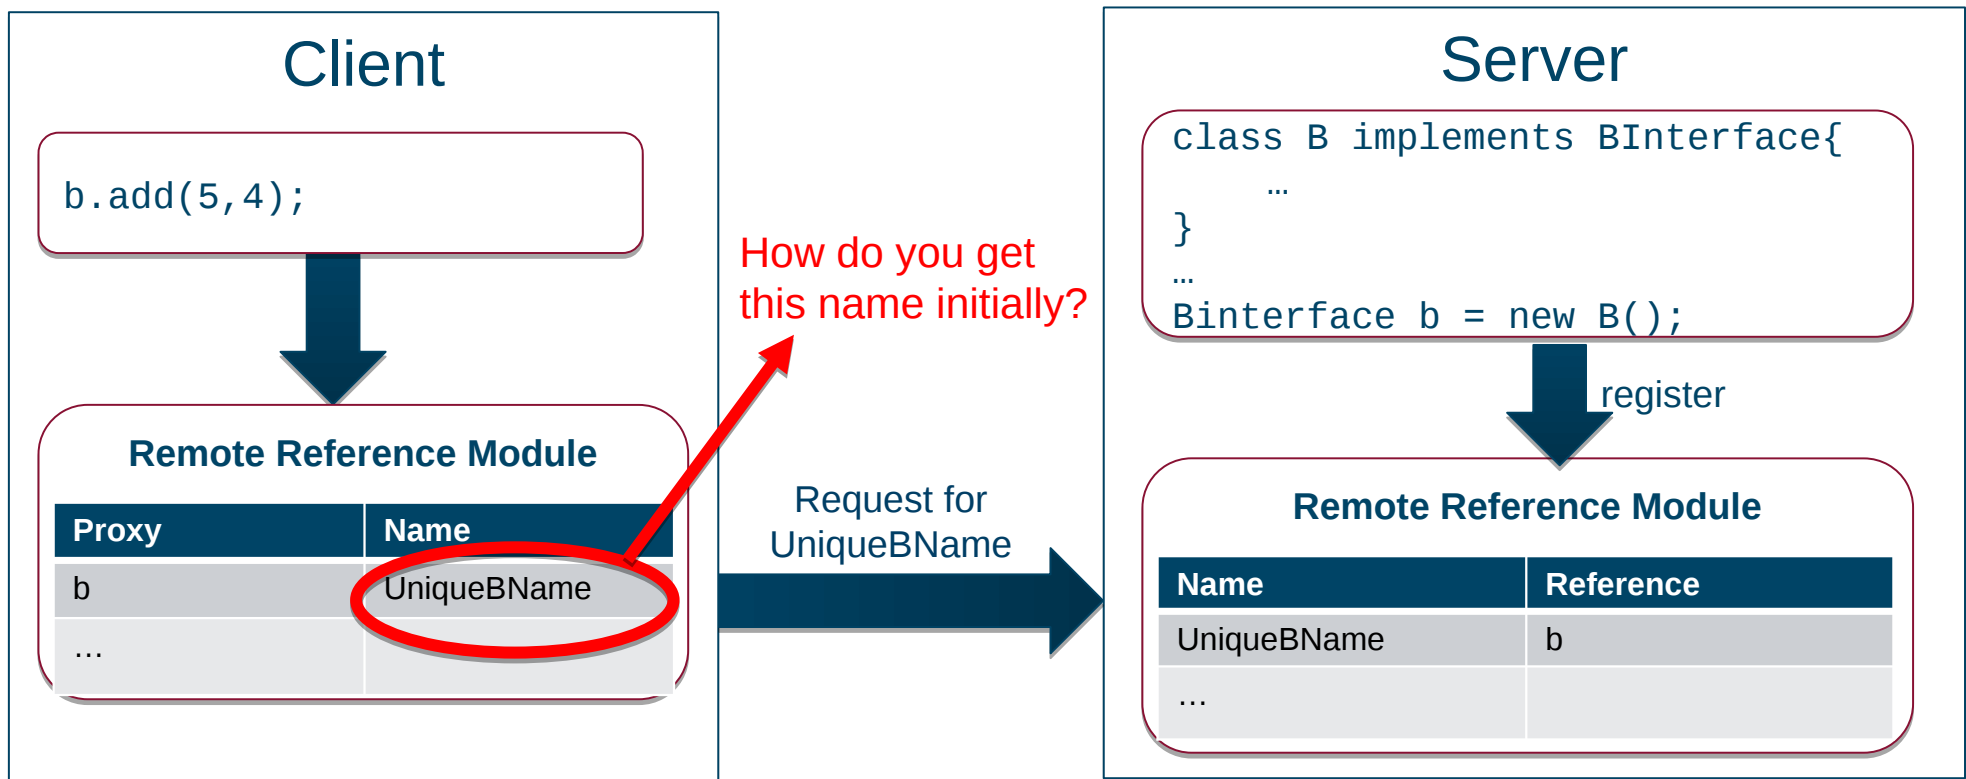


# Remote Reference Module

Is responsible for translating between local and remote object references and for creating remote object references.

Has a remote object table

- **Server-side:** An entry for each remote object held by any process. E.g., B at host B.
- **Client-side:** An entry for each local proxy. E.g., proxy-B at host A



# Binding Service

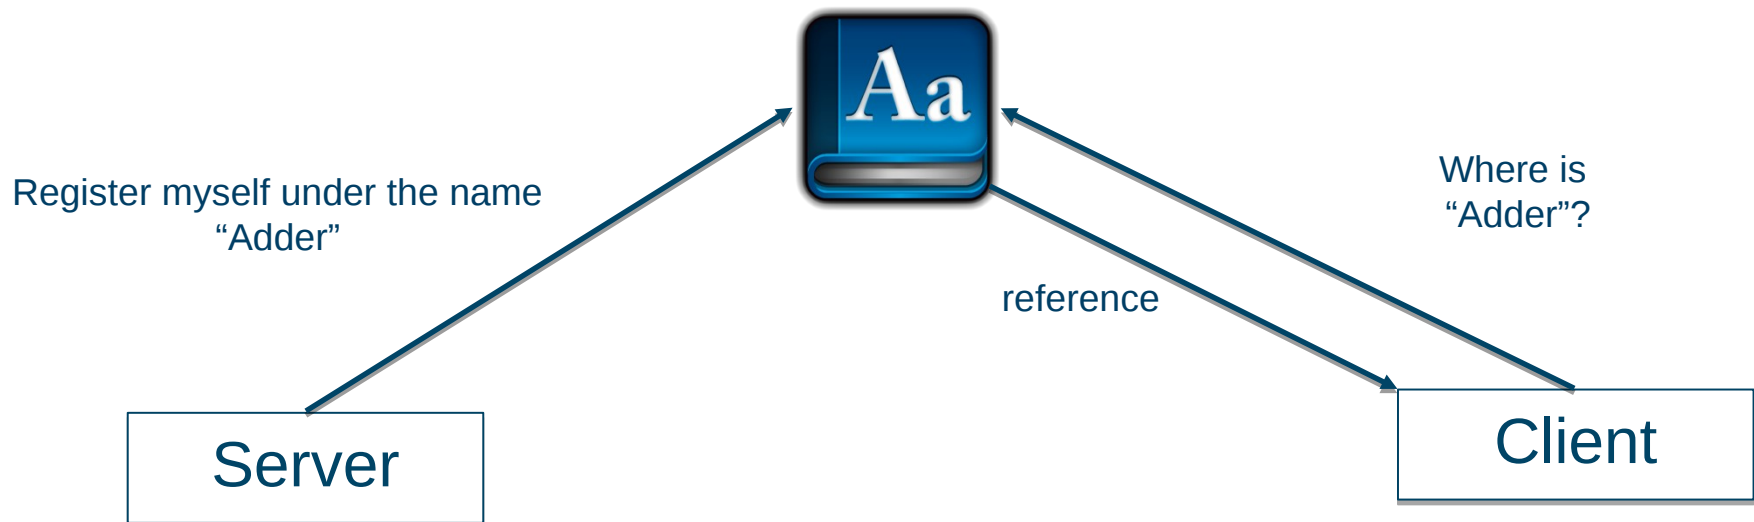
*How do you get the initial reference?*

## **Option 1) Contact the server**

- Go to the server and ask for the name
- What if multiple processes are running of the same class?

## **Option 2) Contact a Binding Service**

- Dictionary service for Middleware
- Or DNS system for RPC
- 



# Do I need to write all this code?



**NO!**

Generated by middleware  
... based on IDL



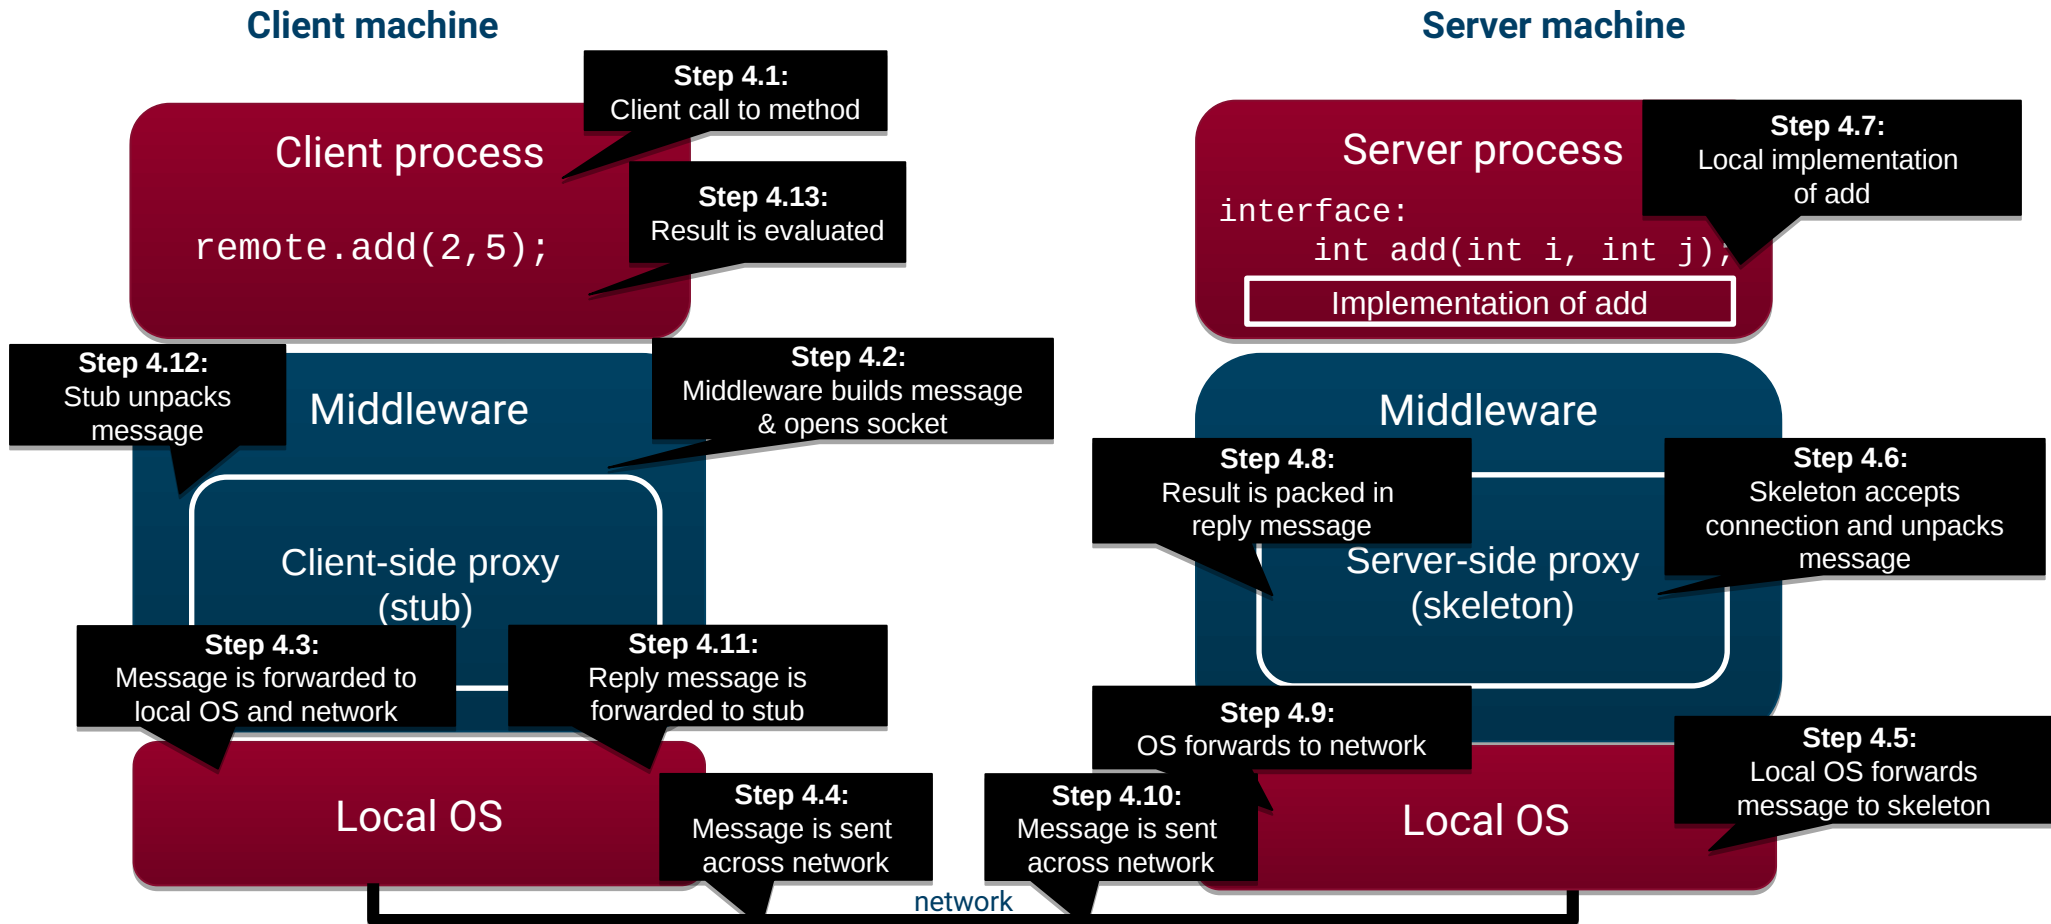
# Putting it all together

## Example: remote method that adds two numbers

STEP 1: Define remote interface, using Interface Definition Language

STEP 2: Implement remote methods

STEP 3: Middleware automatically generates proxies



# Questions ?



# Middleware & Communication