


Distributed Storage


[Hadoop & Friends]

Agenda for today



Distributed storage

- Moving beyond classical storage
- HDFS as a use case



Distributed processing

- MapReduce paradigm
- Improvements via Spark

Algorithmic complexity

Bubble sort, quick sort, Radix sort...

Which operations should we count?

- multiply / divide
- add / subtract

Nope!

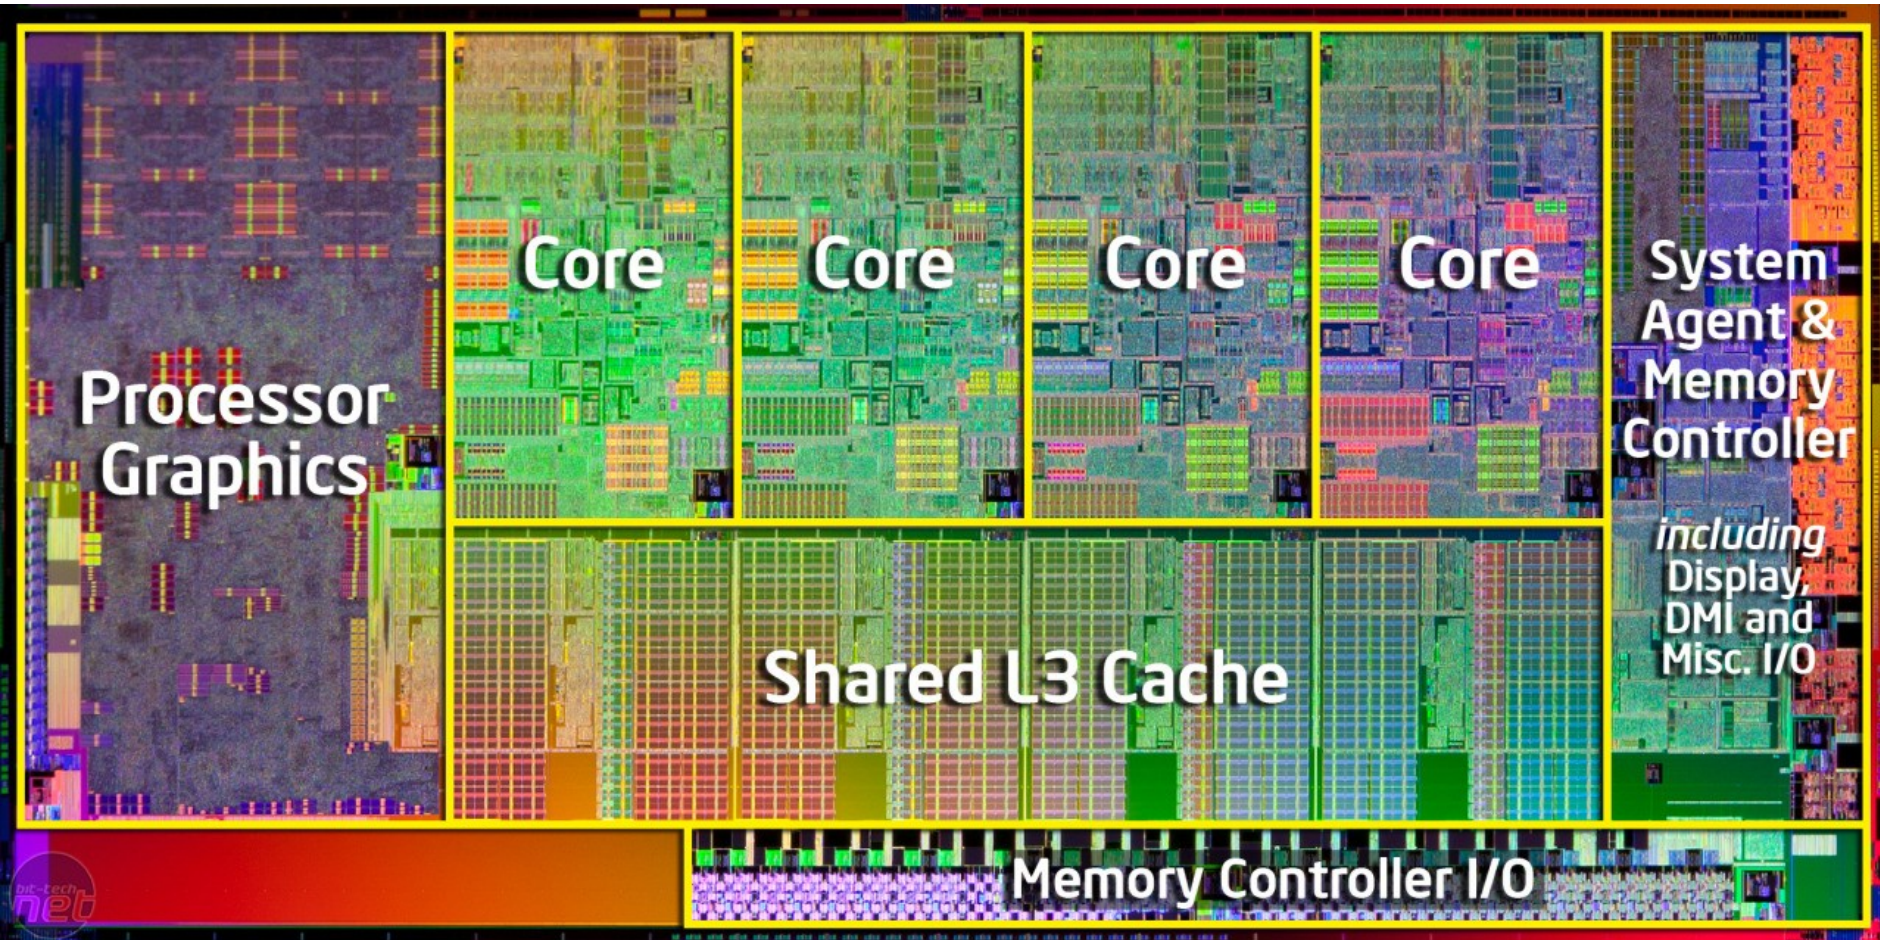
- page fault
- cache miss
- memory access
- disk access (swap space)
- network fetch
- communication between workers

$$O(n^2)$$


$$O(N^{2.8074})$$

$$O(n \log n)$$

Paradigm shift



Moore's law



SOURCE: RAY KURZWEIL, "THE SINGULARITY IS NEAR: WHEN HUMANS TRANSCEND BIOLOGY", P.67, THE VIKING PRESS, 2006. DATAPPOINTS BETWEEN 2000 AND 2012 REPRESENT BCA ESTIMATES.

Interpretation of Moore's law

Until 2005 faster execution

Since 2005 parallel execution


Why?

speed of light

atomic boundaries

limited 3D layering

⇒ **Paradigm Shift**



Basics of parallel processing

[Introduction to Hadoop]

Basics of parallel processing

[Introduction to Hadoop]

What is Hadoop?



Apache top level project, open-source implementation of **frameworks for reliable, scalable, distributed computing and data storage.**

It is a flexible and highly-available architecture for large scale computation and data processing **on a network of commodity hardware.**

What happens in a Google Cluster?

- 1000 individual machine failures
- 1000's of disk failures
- 1 PDU failure (~500-1000 machines disappear for ~6 h)
- 20 rack failures (40-80 machines disappear for 1-6 hours)
- 5 racks go wonky (40-80 machines see 50% packet loss)
- 3 router failures (have to immediately pull traffic for 1h)
- ...

Most (large) jobs see failures!

No, we're not smart enough to program around it

Google Origins

2003

The Google File System

Sanjay Ghemawat, Howard Gobioff, and Shun-Tak Leung
Google*



2004

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

Google, Inc.



2006

Bigtable: A Distributed Storage System for Structured Data

Fay Chang, Jeffrey Dean, Sanjay Ghemawat, Wilson C. Hsieh, Deborah A. Wallach
Mike Burrows, Tushar Chandra, Andrew Fikes, Robert E. Gruber

{fay,jeff,sanjay,wilson,hkerr,m3b,tushar,fikes,gruber}@google.com

Google, Inc.



Google MapReduce


A framework for processing LOADS of data

Framework's job: fault tolerance, scaling & coordination

Programmer's job: write program in MapReduce form



Hadoop is...



HDFS


Hadoop Distributed File System

**Big Data
Storage**

&



**Big Data
Processing**




HDFS

Hadoop Distributed File System

HDFS: storing large files

300 MB



Four score and seven years ago our fathers brought forth on this continent, a new nation, conceived in Liberty, and dedicated to the proposition that all men are created equal.

Now we are engaged in a great civil war, testing whether that nation, or any nation so conceived and so dedicated, can long endure. We are met on a great battle-field of that war.

We have come to dedicate a portion of that field, as a final resting place for those who here gave their lives that that nation might live. It is altogether fitting and proper that we...

We have a file

HDFS: storing large files

128 MB {

Four score and seven years ago our fathers brought forth on this continent, a new nation, conceived in Liberty, and dedicated to the proposition that all men are created equal.

128 MB {

Now we are engaged in a great civil war, testing whether that nation, or any nation so conceived and so dedicated, can long endure. We are met on a great battle-field of that war.

44 MB {

We have come to dedicate a portion of that field, as a final resting place for those who here gave their lives that that nation might live. It is altogether fitting and proper that we...

HDFS splits it
into blocks

HDFS: storing large files

x3

Four score and seven
years ago our fathers
brought forth on this
continent, a new nation,
conceived in Liberty,
and dedicated to the
proposition that all men
are created equal.

x3

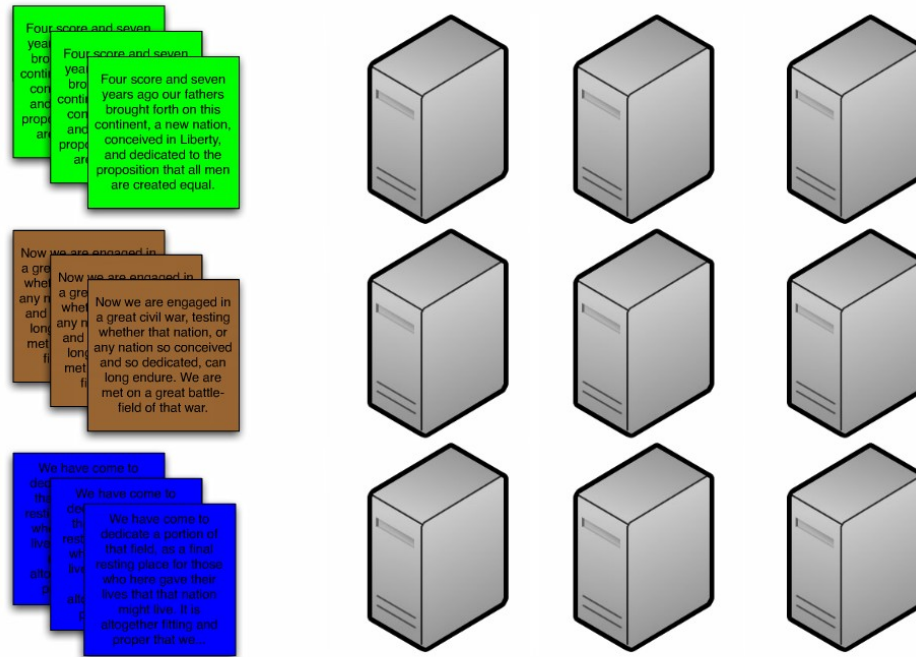
Now we are engaged in
a great civil war, testing
whether that nation, or
any nation so
conceived and so
dedicated, can long
endure. We are met on
a great battle-field of
that war.

x3

We have come to
dedicate a portion of
that field, as a final
resting place for those
who here gave their
lives that that nation
might live. It is
altogether fitting and
proper that we...


HDFS will keep 3
copies of each block

HDFS: storing large files




HDFS stores these
blocks on datanodes

HDFS: storing large files




HDFS distributes the
blocks to the DN's

HDFS: storing large files




The NameNode tracks
blocks and Datanodes

HDFS: storing large files



Sometimes a Datanode
will die. Not a problem.

HDFS: storing large files



NameNode tells other
datanodes to copy blocks,
back to 3x replication



Hadoop MapReduce

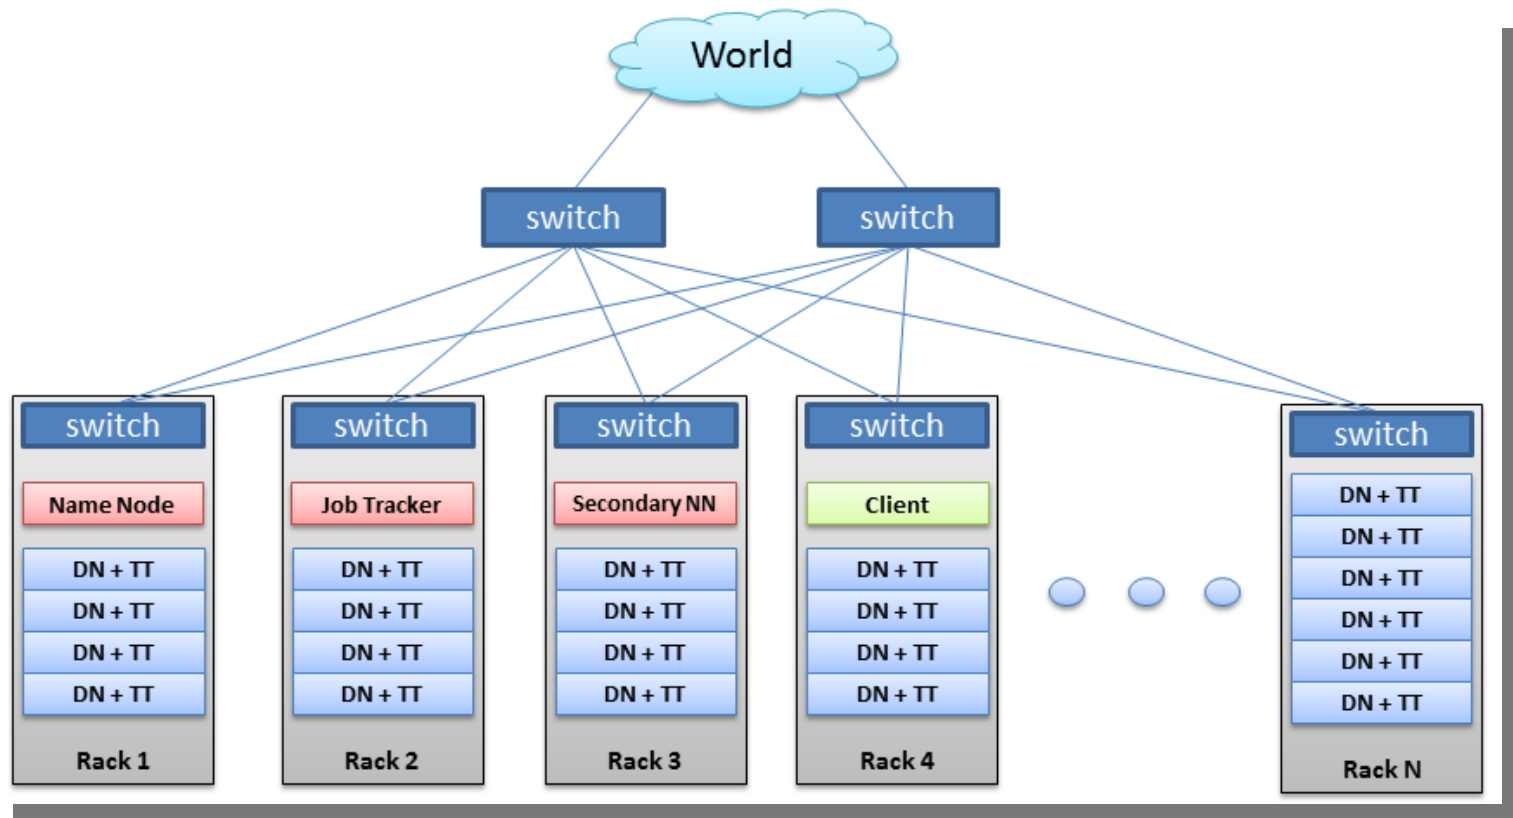
Model for distributed processing

Mantra:

Move the computation to the data

In practice this means running the processing on a machine whose HDFS Datanode holds the data

Hadoop architecture



TT : task tracker

Hadoop MapReduce

Data in the form of a <key, value> pair

- § <byte, text>
- § <user id, user profile>
- § <timestamp, log entry>
- § <user id, list of user id's of friends>
- § ...


Inspired by list processing (Lisp) , functional programming:

- § immutable data
- § pure functions (no side effects): map, reduce

Simple model = easy to reason about

Google MapReduce

Map: map each $\langle \text{key}, \text{value} \rangle$ of input list onto 0, 1, or more pairs of type $\langle \text{key}_2, \text{value}_2 \rangle$ of output list




Behavior:

- Map to 0 elements in the output → filtering
- Map to +1 elements in the output → distribution

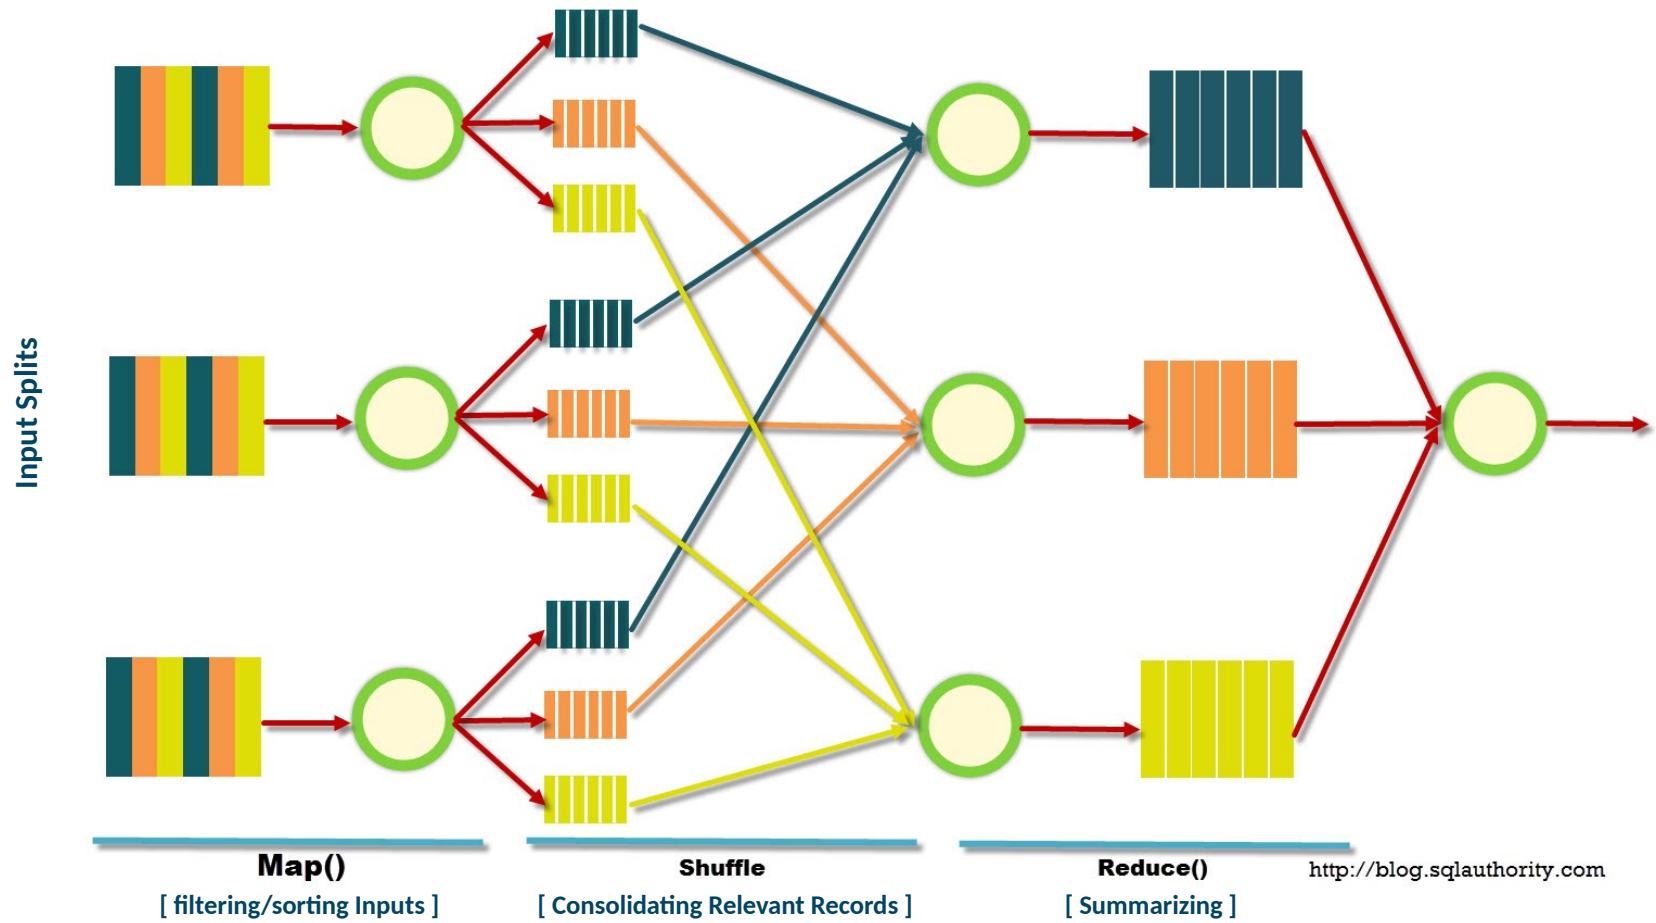
Google MapReduce

Reduce: combine the $\langle \text{key}, \text{value} \rangle$ pairs of the input list to an aggregate output value



Bringing it together

How MapReduce Works?



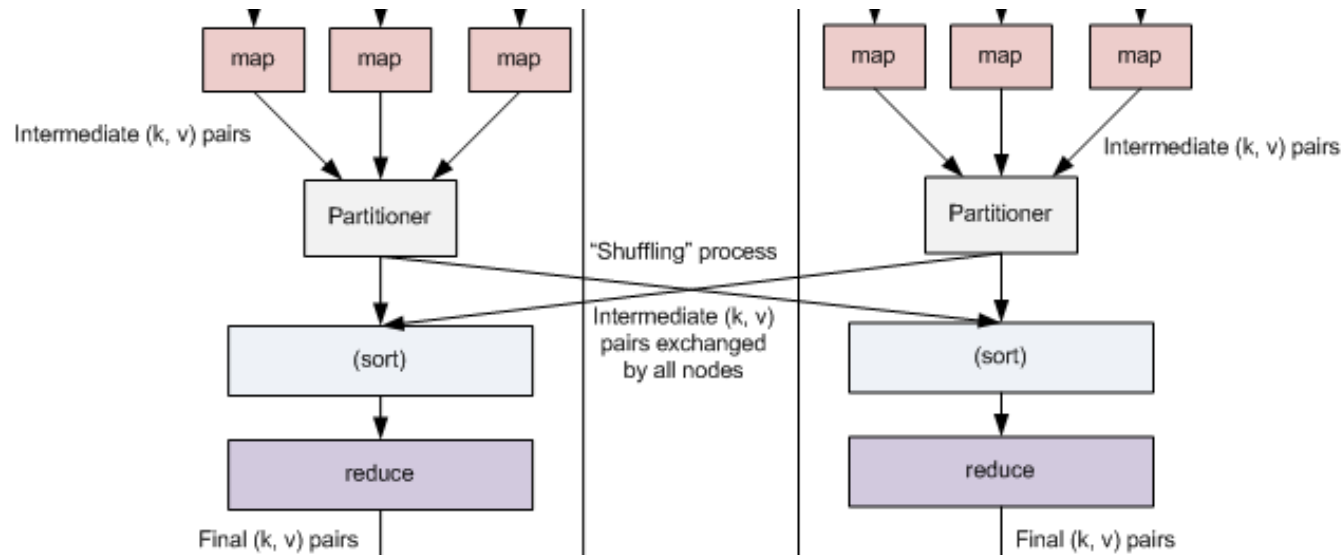
What you need to write

Mapper: application code

Partitioner: send data to correct Reducer machine

Sort: group input from different mappers by key

Reducer: application code



TokenizerMapper

```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

TokenizerMapper

```
public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

    private final static IntWritable one = new IntWritable(1);
    private Text word = new Text();

    public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {

        StringTokenizer itr = new StringTokenizer(value.toString());
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            context.write(word, one);
        }
    }
}
```

Breaking value
into tokens

Pushing every token into context

IntSumReducer

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

IntSumReducer

```
public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
    private IntWritable result = new IntWritable();

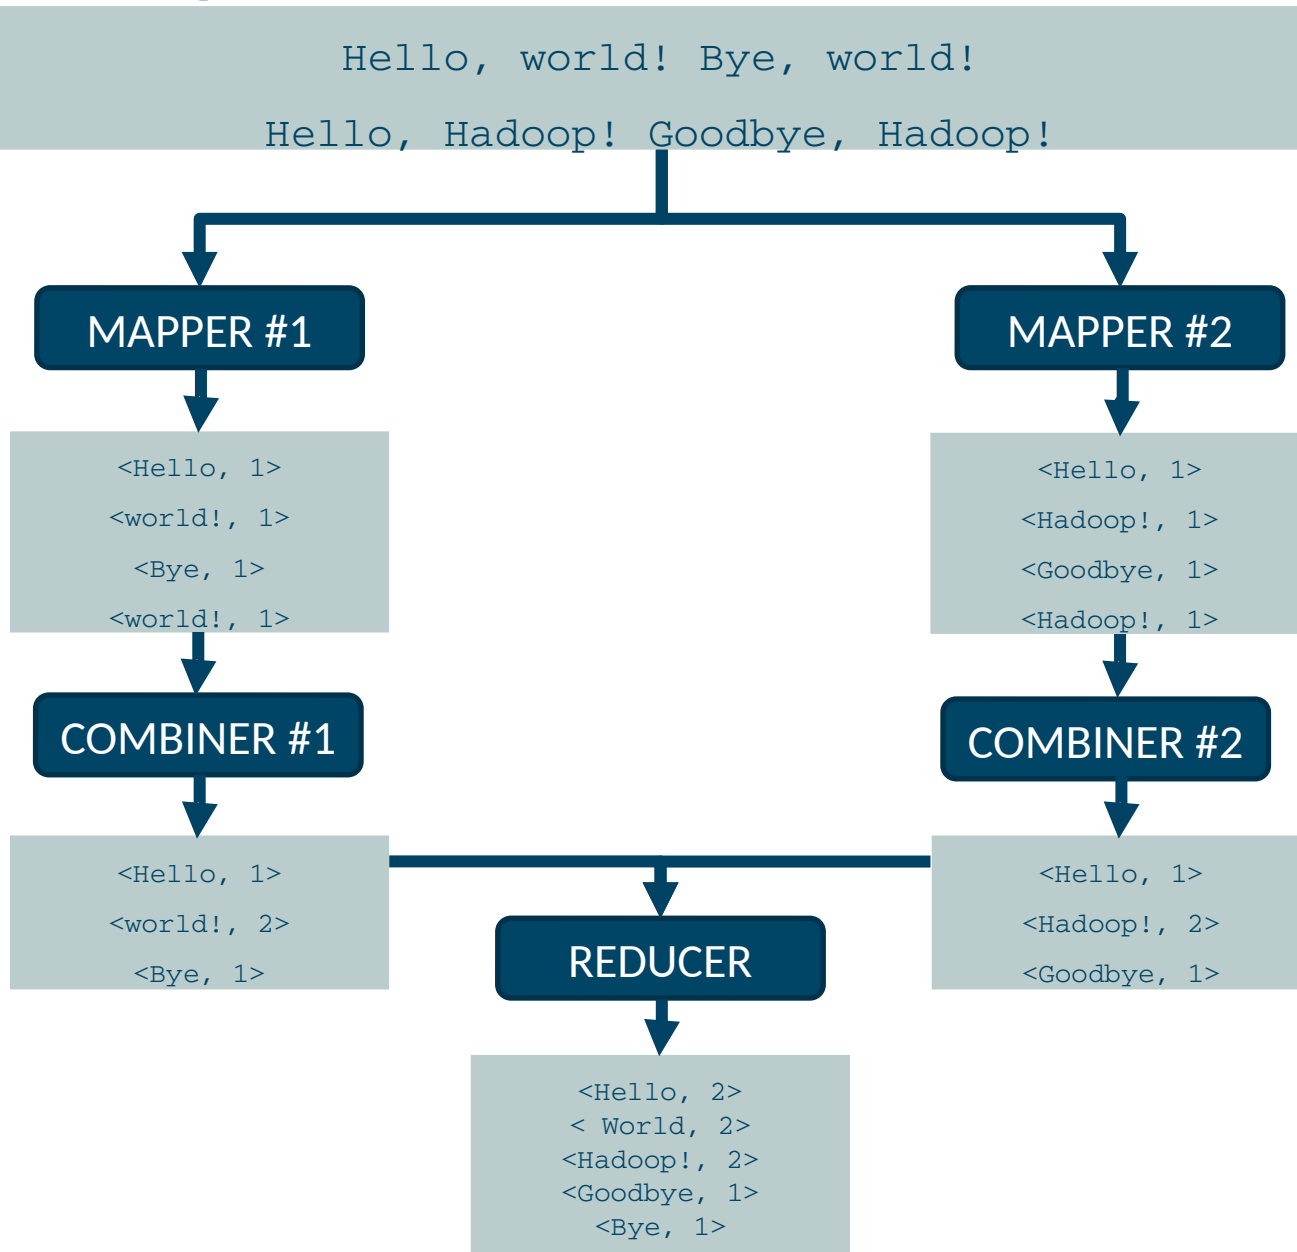
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values) {
            sum += val.get();
        }
        result.set(sum);
        context.write(key, result);
    }
}
```

Accumulating all the values (int)

Pushing the sum (result) into context

Small example



Hadoop Architecture Overview




Cluster

Components

- **MapReduce Framework:** implement MapReduce paradigm
- **Cluster:** host machines (nodes).
- **HDFS federation:** provides logical distributed storage.
- **YARN Infrastructure:** assign resources (CPU, memory, etc.=

Hadoop Architecture Overview



Components

- **MapReduce Framework:** implement MapReduce paradigm
- **Cluster:** host machines (nodes).
- **HDFS federation:** provides logical distributed storage.
- **YARN Infrastructure:** assign resources (CPU, memory, etc.=

YARN Infrastructure

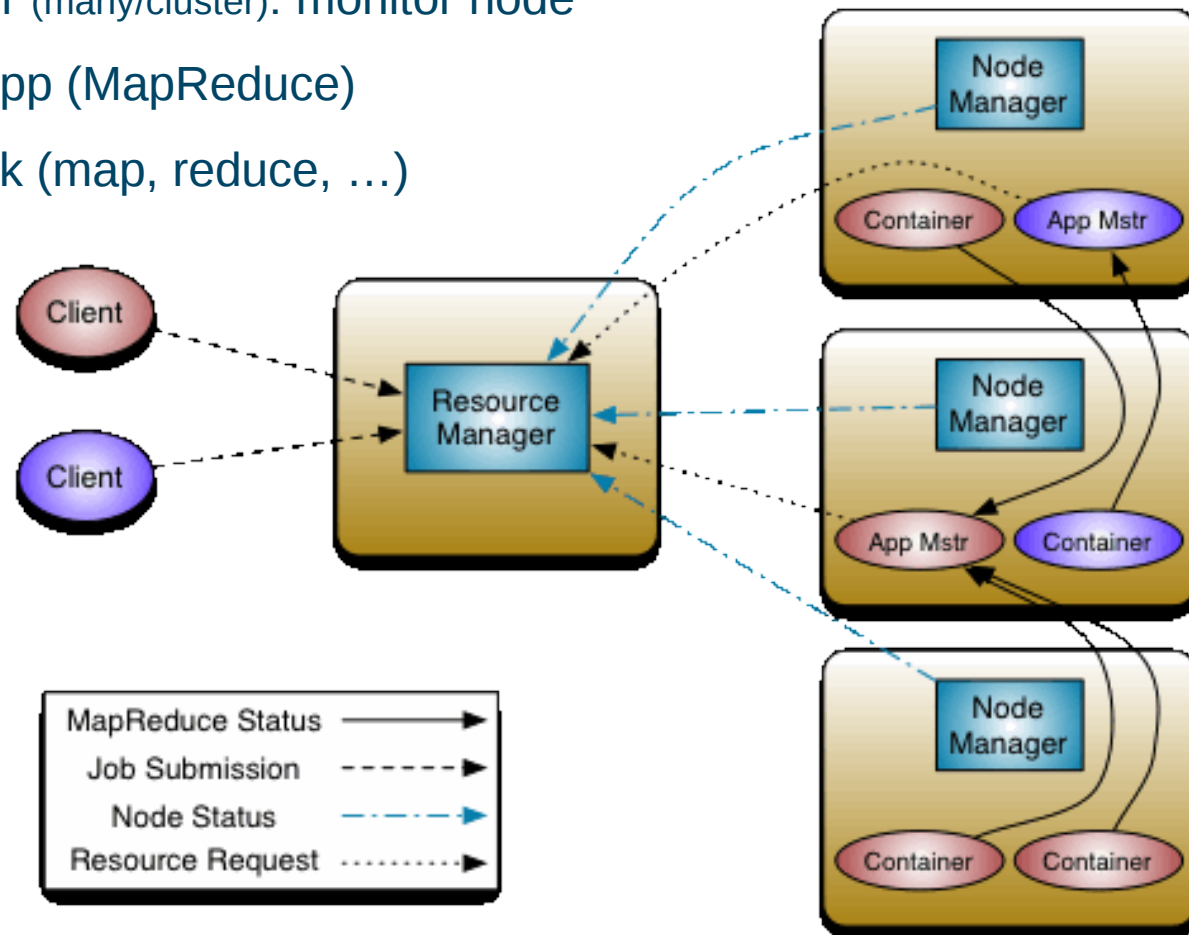
Yet Another Resource Negotiator

Resource Manager (1/cluster): assign cluster resources to applications

Node Manager (many/cluster): monitor node

App Master: app (MapReduce)

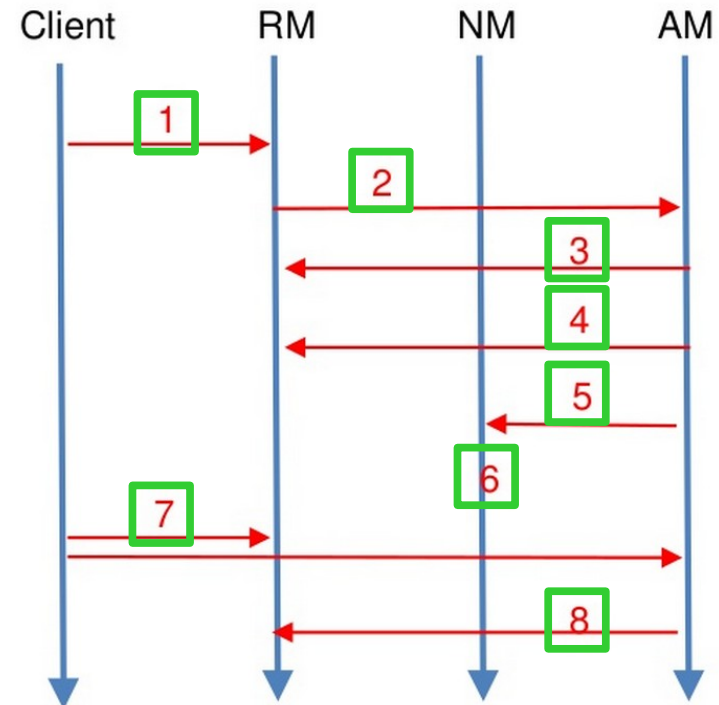
Container: task (map, reduce, ...)



YARN application lifecycle

RM: resource manager
NM: node Manager
AM: application master

1. Client submits app
2. RM allocates AM container
3. AM registers with NM
4. AM requests containers from RM
5. AM tells NM to launch containers
6. Application code is executed
7. Monitor app status in RM/AM
8. AM unregisters with RM



© Hortonworks

Application developers only need to write code for 6

Break

[See you in 15 mins.]


Hadoop in the wild

Hadoop in action: real-life examples

Largest Hadoop cluster in Europe?: 1650 machines (~20K jobs/day)

Use:

- Reporting to record labels so everyone gets paid
- Creating top lists of what is the most popular music right now
- Getting feedback on different aspects of the product
 - improve user experience
- Powering intelligent radio and discovery features



Data Infrastructure:

- 1650 Hadoop Nodes
- 65 PB Storage, 70 TB RAM
- 20 TB data ingested via Kafka/day
- 200 TB generated by Hadoop/day

Hadoop in action: real-life examples



What is trending? What causes a tweet to be trending?
General statistics, etc.



Ad reporting. Fake news detection. Recommendation, etc.



Crawling the web



Will you switch banks?
Linked with social media data



Web clicks, where are users coming from?
Following up on user visits

Solr 

 STORM

APACHE
HBASE


APACHE
Spark 



OOZIE



 kafka




 mahout

the Hadoop eco-system

Pig: scripting language

```
a = load '/user/porky/word_count_text.txt';  
b = foreach a generate flatten(TOKENIZE((chararray)  
$0)) as word;  
c = group b by word;  
d = foreach c generate COUNT(b), group;  
store d into '/user/porky/pig_wordcount';
```



Hive: SQL language

```
CREATE TABLE docs (line STRING);  
LOAD DATA INPATH 'text' OVERWRITE INTO TABLE docs;  
CREATE TABLE word_counts AS  
SELECT word, count(1) AS count FROM  
(SELECT explode(split(line, '\s')) AS word FROM docs) word  
GROUP BY word  
ORDER BY word;
```



Kafka: the Log

Distributed commit log / message service

Store all events: richer than a DB (= only last value)

Will be HUGE once the Internet of Things matures



Multithreaded programming



Multithreaded programming



Shortcoming of MapReduce

Forces your data processing into **MAP** and **REDUCE**


- Other workflows missing include join, filter, flatMap, groupByKey, union, intersection, ...

Based on “Acyclic Data Flow” from Disk to Disk (HDFS)

- Not efficient for iterative tasks, i.e. Machine Learning

Only for Batch processing

- Interactivity, streaming data



Hadoop and disks



Hard drive access is killing performance and blocking functionality

One Solution is Apache Spark

Works on top of Hadoop, HDFS,


Has many other workflows, i.e. join, filter, flatMapdistinct, groupByKey, reduceByKey, sortByKey, collect, count, first...
(around 30 efficient **distributed operations**)

In-memory caching of data (for iterative, graph, and machine learning algorithms, etc.)




Spark Uses Memory instead of Disk

Hadoop: Use Disk for Data Sharing




Spark Uses Memory instead of Disk

Hadoop: Use Disk for Data Sharing



Spark: In-Memory Data Sharing




Sort competition

	Hadoop MR Record (2013)	Spark Record (2014)	Spark, 3x faster with 1/10 the nodes
Data Size	102.5 TB	100 TB	
Elapsed Time	72 mins	23 mins	
# Nodes	2100	206	
# Cores	50400 physical	6592 virtualized	
Cluster disk throughput	3150 GB/s (est.)	618 GB/s	
Network	dedicated data center, 10Gbps	virtualized (EC2) 10Gbps network	
Sort rate	1.42 TB/min	4.27 TB/min	
Sort rate/node	0.67 GB/min	20.7 GB/min	

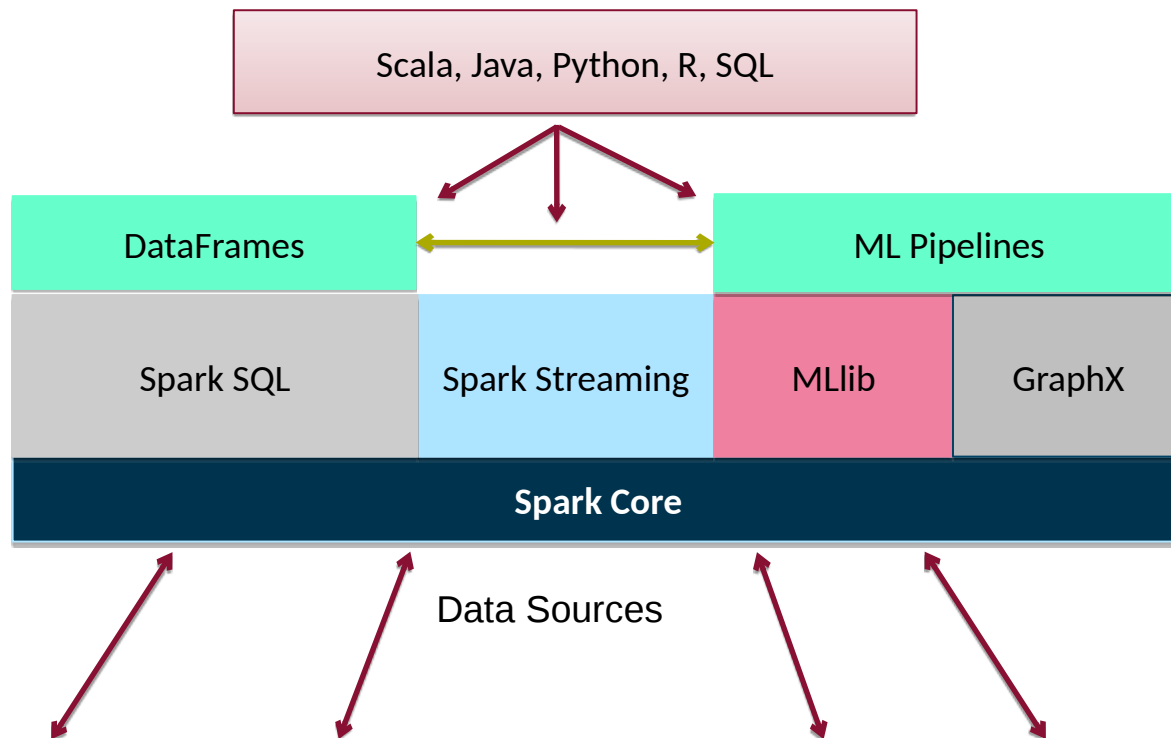
Sort benchmark, Daytona Gray: sort of 100 TB of data (1 trillion records)
<http://databricks.com/blog/2014/11/05/spark-officially-sets-a-new-record-in-large-scale-sorting.html>

Logistic regression performance



Apache Spark

Apache Spark supports data analysis, machine learning, graphs, streaming data, etc. It can read/write from a **range of data types** and allows **development in multiple languages**.



Hadoop HDFS, HBase, Hive, Apache S3, Streaming, JSON, MySQL, and HPC-style (GlusterFS, Lustre)

Resilient Distributed Datasets (RDDs)

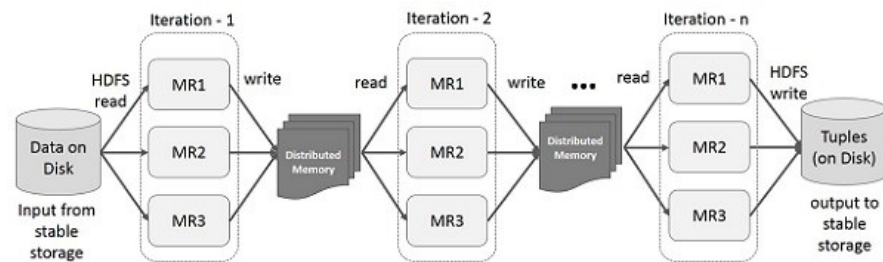

Immutable distributed collection of objects

All Spark components use RDDs

Use transformations to create new RDDs

- From storage
- From other RDDs

Fault tolerant




DataFrames & SparkSQL

Organize the data in named columns

Similar to a relational database...

- Immutable once constructed
- Enable distributed computations



DataFrames & SparkSQL


Organize the data in named columns

Similar to a relational database...

- Immutable once constructed
- Enable distributed computations

How to construct Dataframes

- Read from file(s)
- Transforming an existing DFs
- Parallelizing a python collection list
- Apply transformations and actions



DataFrame example

```
// Create a new DataFrame that contains "students"  
students = users.filter(users.age < 21)
```

```
//Count the number of students users by gender  
students.groupBy("gender").count()
```

```
// Join young students with another DataFrame  
called logs  
students.join(logs, logs.userId == users.userId,  
"left_outer")
```

RDDs vs. DataFrames


RDDs **provide a low level interface** into Spark

DataFrames **have a schema**

DataFrames are **cached and optimized by Spark**

DataFrames are **built on top of the RDDs and the core Spark API**


Example: performance



Spark Operations

Transformations (create a new RDD)	map filter sample groupByKey reduceByKey sortByKey intersection	flatMap union join cogroup cross mapValues reduceByKey
Actions (return results to driver program)	collect Reduce Count takeSample take lookupKey	first take takeOrdered countByKey save foreach

Directed Acyclic Graphs (DAG)



DAGs track dependencies
(also known as Lineage)

- nodes are RDDs
- arrows are Transformations

Why?

- *Program resonates with humans and computers*
- Improvement via:
 - Sequential access to data
 - Predictive processing

Narrow Vs. Wide transformation

Narrow

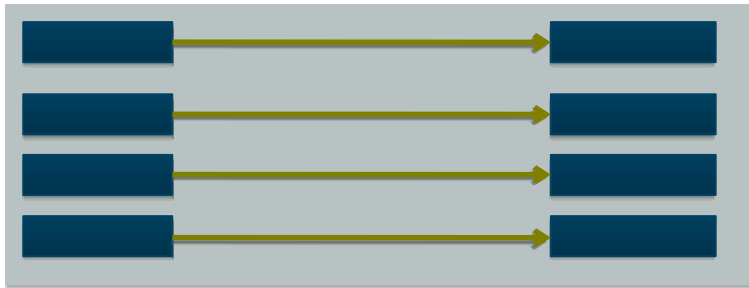


Map

Required elements for computation in a single partition **live in the single partition** of parent RDD

Narrow Vs. Wide transformation

Narrow





Map

Required elements for computation in a single partition **live in the single partition** of parent RDD

Vs.

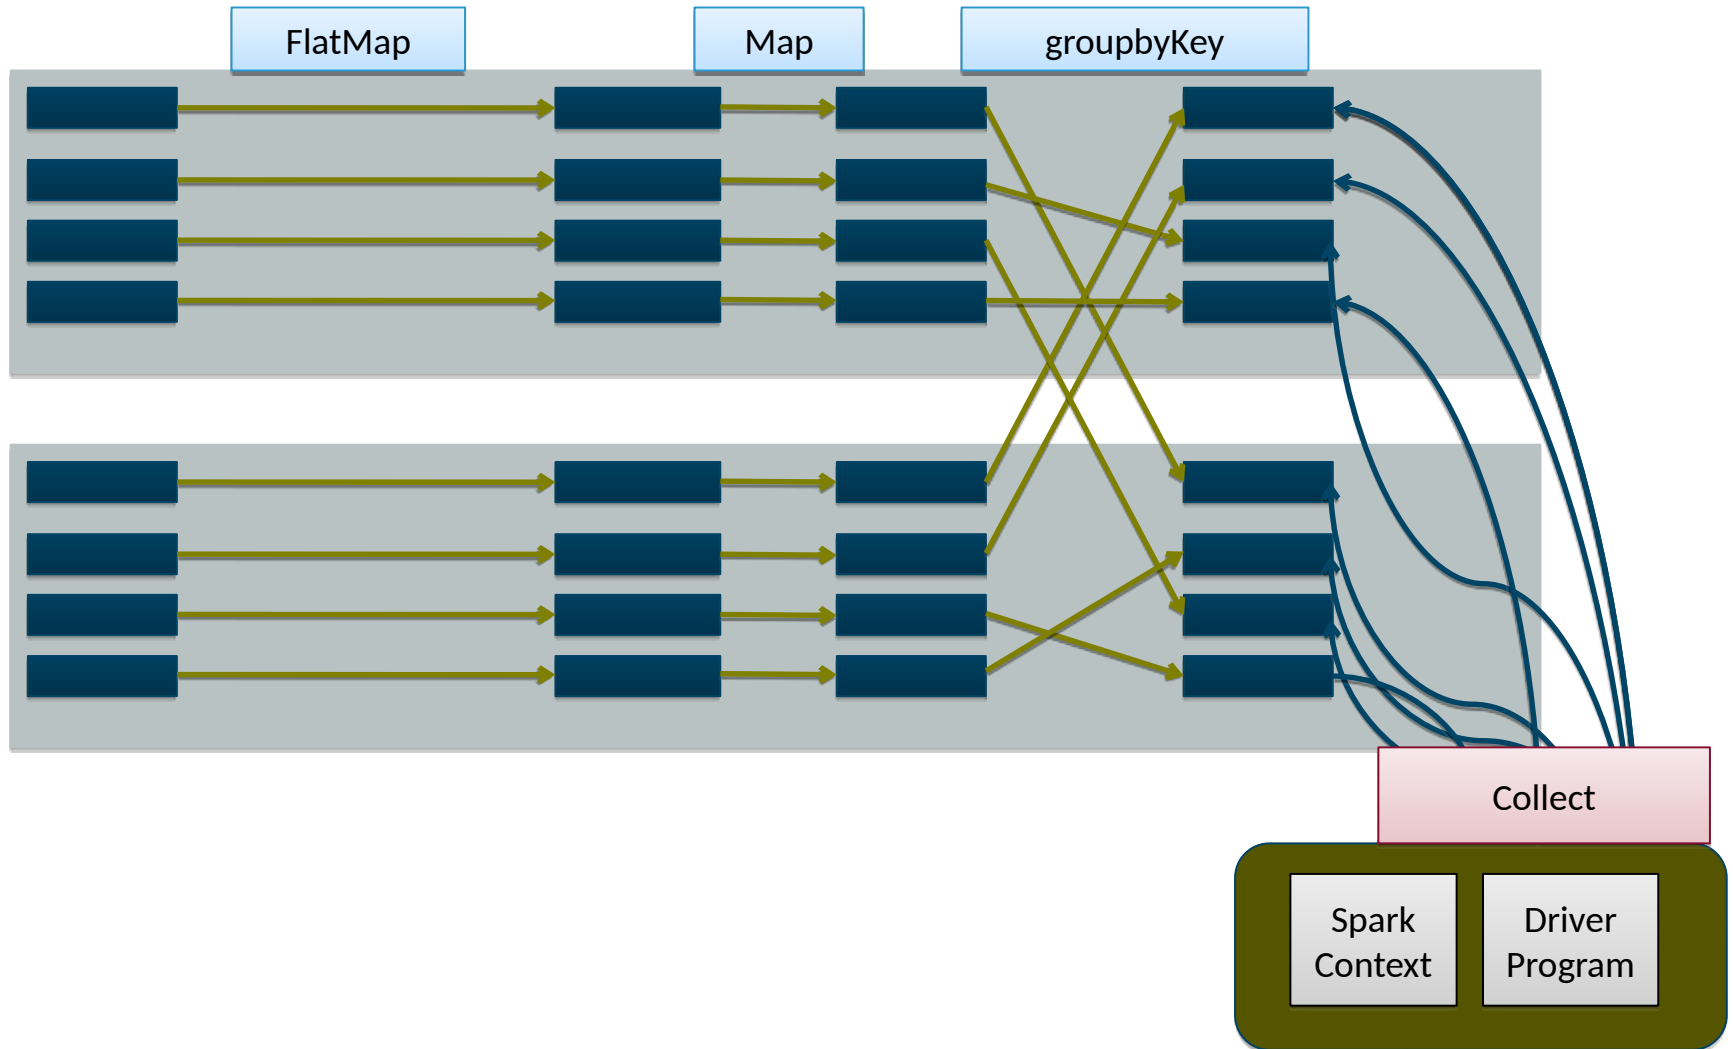
Wide



groupByKey

Required elements for computation in a single partition **may live in many partitions** of parent RDD

Spark Workflow



Python RDD API Examples

Word count

```
text_file = sc.textFile("hdfs://usr/godil/text/book.txt")
counts = text_file.flatMap(lambda line: line.split(" ")) \
    .map(lambda word: (word, 1)) \
    .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://usr/godil/output/wordCount.txt")
```

Logistic Regression

```
# Every record of this DataFrame contains the label and
# features represented by a vector.
df = sqlContext.createDataFrame(data, ["label", "features"])
# Set parameters for the algorithm.
# Here, we limit the number of iterations to 10.
lr = LogisticRegression(maxIter=10)
# Fit the model to the data.
model = lr.fit(df)
# Given a dataset, predict each point's label, and show the results.
model.transform(df).show()
```

Spark's Main Use Cases

Streaming Data

Machine Learning

Interactive Analysis

Data Warehousing

Batch Processing

Exploratory Data Analysis

Graph Data Analysis

Spatial (GIS) Data Analysis

And many more

Spark in the Real World (I)



Uber – the online taxi company gathers terabytes of event data from its mobile users every day.

- By using Kafka, Spark Streaming, and HDFS, to build a continuous ETL (*extract, transform, load*) pipeline
- Convert raw unstructured event data into structured data as it is collected
- Uses it further for more complex analytics and optimization of operations



Pinterest – Uses a Spark ETL pipeline

- Leverages Spark Streaming to gain immediate insight into how users all over the world are engaging with Pins—in real time.
- Can make more relevant recommendations as people navigate the site
- Recommends related Pins
- Determine which products to buy, or destinations to visit

Spark: when not to use

Even though Spark is versatile, that doesn't mean Spark's in-memory capabilities are the best fit for all use cases:

- For many **simple use cases** Apache MapReduce and Hive might be a more appropriate choice
- Spark was not designed as a **multi-user environment**
- Spark users are **required to know that memory they have is sufficient** for a dataset
- **Adding more users adds complications**, since the users will have to coordinate memory usage to run code

Hadoop Ecosystem



APACHE
DRILL

Interactive Analysis

Stream Processing



Apache Storm

Data Transfer



Data Streaming
(Unstructured)



ZOO KEEPER
Coordination Service



PIG

Scripting
Language



Machine
Learning



HiveQL Query



Column Datastore

APACHE
HBASE




Core Hadoop

hadoop




HDFS Core Hadoop

Summarizing



Distributed storage

- HDFS and related technologies enable resilient storage



Distributed processing

- MapReduce paradigm
- Spark enables interactive processing and a richer set of operations.

Questions?

Further Reading

- **Hadoop**

<http://hadoop.apache.org/>

<https://cwiki.apache.org/confluence/display/HADOOP2>

<https://ercoppa.github.io/HadoopInternals/HadoopArchitectureOverview.html>

- **MapReduce**

<https://blog.sqlauthority.com/2013/10/09/big-data-buzz-words-what-is-mapreduce-day-7-of-21/>

<https://hadoop.apache.org/docs/current/api/org/apache/hadoop/mapreduce/Mapper.html>

<https://hadoop.apache.org/docs/current/api/org/apache/hadoop/mapreduce/Reducer.html>

<https://hadoop.apache.org/docs/current/api/org/apache/hadoop/mapreduce/Partitioner.html>

- **Scaling Computing @Spotify**

<https://www.youtube.com/watch?v=cdsfRXr9pJU>

<https://www.slideshare.net/RafaWojdya/the-evolution-of-hadoop-at-spotify-through-failures-and-pain>

<https://www.slideshare.net/JoshBaer/how-apache-drives-music-recommendations-at-spotify>

- **Evolution and Limits of Computation**

Markov, I. Limits on fundamental limits to computation. Nature 512, 147–154 (2014).

Published Article: <https://doi.org/10.1038/nature13570> arXiv:1408.3821.

Preprint: <https://arxiv.org/abs/1408.3821> (*free access*)

Distributed Storage

[Hadoop & Friends]