



University
of Antwerp

Specification and Verification

Lecture 1: Introduction - To verify or not to verify?

Guillermo A. Pérez

September 27, 2024

Outline

- 1 Course particularities
- 2 Why should we worry?
- 3 Formal verification in a nutshell
- 4 Conclusion & outline

TL;DR: This course in short

What is Verification?

Formal verification is the act of proving or disproving the correctness of intended algorithms underlying a system with respect to a certain formal specification or property, using formal methods of mathematics.

Why verify?

Essentially: To avoid money and life loss.

References

Main references

- Christel Baier, Joost-Pieter Katoen: **Principles of Model Checking**. MIT Press 2018.
- Mickael Randour: Verification course @ UMONS.
- **Lectures in Game Theory for Computer Scientists**. Eds: K. Apt, E. Grädel. Cambridge University Press 2011.
- **Handbook of Model Checking**. Eds: E.M. Clarke, T.A. Henzinger, H. Veith, R. Bloem. Springer 2018.

Required and target competences

What tools do we need?

Discrete maths, Modelling, Algorithms & complexity, Formal language theory

What skills will we obtain?

We will get an overview of the verification field and

- Prepare a symbolic solver, split into two deliveries:
 1. Library with graph/automata algorithms using BDDs
 2. Parity-game solver

How will these skills be useful?

To assert strong correctness guarantees about your code/hardware.

Evaluation method

Project: a verification tool

A project that accounts for **40% of your grade**:

- You prepare a model checker that can deal with provided benchmarks

Exercises: optional, but...

The exercises can be counted towards your exam grade!

Written examination

The **remaining 60%**:

- You can bring **one cheat sheet**
- Answers (report in **L^AT_EX**) for exercise sessions count here! To be turned in before exam date

Outline

- 1 Course particularities
- 2 Why should we worry?
- 3 Formal verification in a nutshell
- 4 Conclusion & outline

Let's talk about bugs...



Insects



Computer bugs

Let's talk about bugs...



Insects



Computer bugs

- Plenty of them

- Plenty of them

Let's talk about bugs...



Insects

- Plenty of them
- Mostly annoying



Computer bugs

- Plenty of them
- Mostly annoying

Let's talk about bugs...



Insects

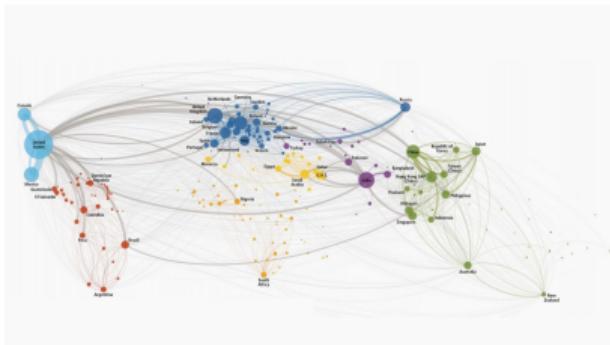
- Plenty of them
- Mostly annoying
- Serve a purpose



Computer bugs

- Plenty of them
- Mostly annoying
- At best useless, can be harmful!

It's all about money (1/4)



AT&T long-distance service (1990)

- Bug caused continuous cascade reboots of all long-distance switches.
- Impact: 9-hour outage.
- Costs: 60-100 million US\$.
- Source: wrong interpretation of C break statement.

It's all about money (2/4)

Pentium FDIV (1994)



- Bug in the floating point **division** unit (FDIV).
- Impact: inaccurate results for $1 \cdot 10^9$ random floating point divisions.
- Costs: ~ 500 million US\$ (replacement of all processors).
 - + PR nightmare for Intel!
- Source: 5 missing entries in a 1066-entry look-up table.

It's all about money (3/4)

Ariane 5 (1996)



- Loss of guidance after 37s followed by self-destruction.
- Costs: > 500 million US\$.
- Source: data conversion from 64-bit floating point to 16-bit signed integer causing overflow in the hardware.
 - Appropriate software handler was **disabled** to improve efficiency.

It's all about money (4/4)



Mars Climate Orbiter (1998)

- Atmosphere entry at wrong angle resulting in disintegration.
- Costs: 327 million US\$ (mission failure).
- Source: ground software sending instructions calculated in the wrong units (imperial vs. metric units as the NASA-Lockheed contract specified).

It's all about safety (1/2)

Therac-25 radiation therapy (1985-1987)

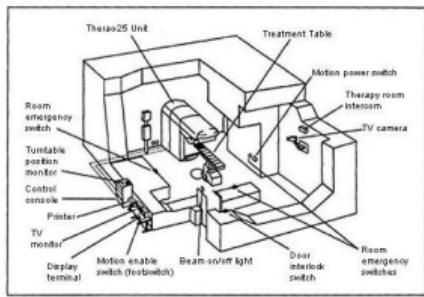
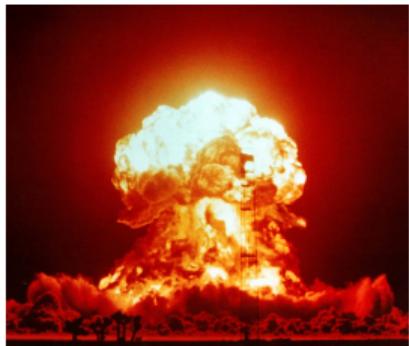


Figure 1. Typical Therac-25 facility

- Two modes: one “safe” direct mode, one powerful mode requiring shielding; bug caused mismatch.
- Impact: **several deaths by radiation poisoning.**
- Source: *race condition* in the software managing the choice of mode.

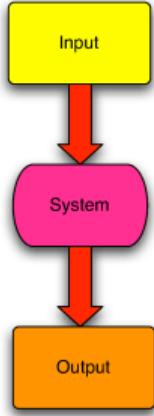
It's all about safety (2/2)



The doomsday bug (1983)

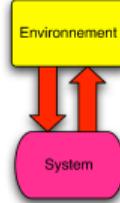
- Soviet nuclear early-warning system (Oko) falsely reports five incoming US missiles.
- Possible impact: WW3?
- Avoided by Stanislav Petrov who judged the report to be a false alarm.
- Source: bug in the Soviet satellite detection system.

Batch processing vs. reactivity



Batch processing system

- Computes results.
- Correctness easier to assess.



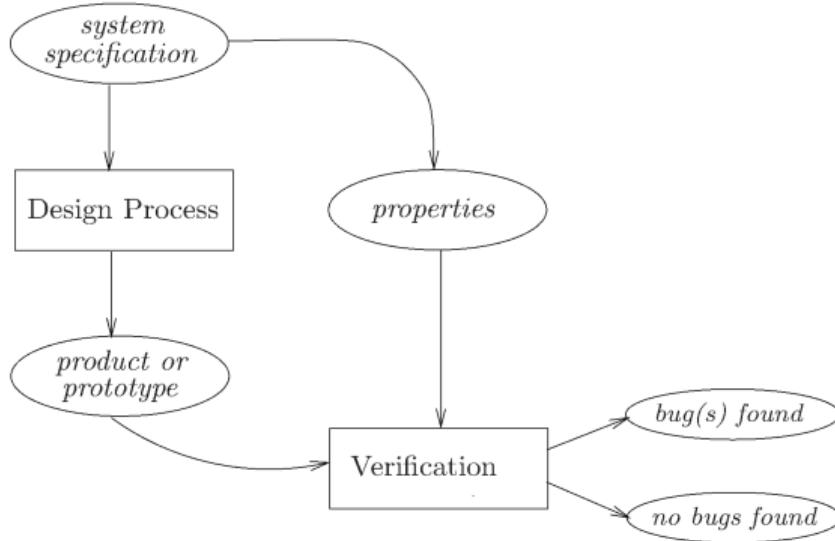
Reactive system

- Continuous interaction with the environment:
 - requests information,
 - reacts to events.
- Correctness very difficult to assess.

Outline

- 1 Course particularities
- 2 Why should we worry?
- 3 Formal verification in a nutshell
- 4 Conclusion & outline

Hardware and software verification (1/4)



A posteriori verification

Properties to check obtained from
the system's **specification**

Hardware and software verification (2/4)

Software verification

- *Peer-reviewing*: static analysis of uncompiled code
 - Useful (catches from 31% to 93%, median 60%, of defects)
 - Used in ~80% of software projects
 - Difficult to catch dynamic issues: concurrency, algorithmic defects...

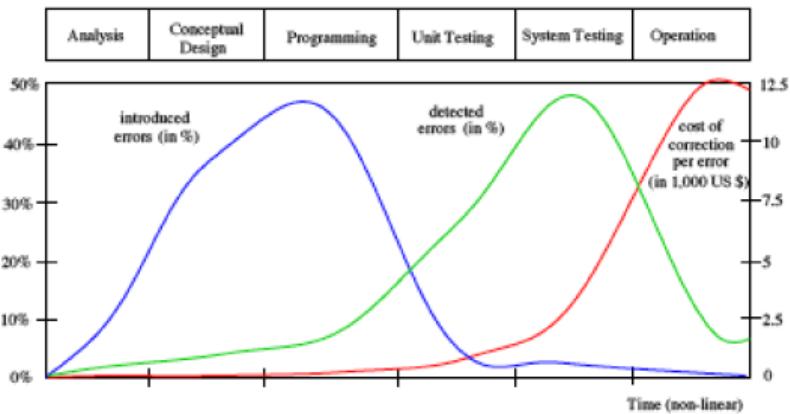
Hardware and software verification (2/4)

Software verification

- *Peer-reviewing*: static analysis of uncompiled code
 - Useful (catches from 31% to 93%, median 60%, of defects)
 - Used in ~80% of software projects
 - Difficult to catch dynamic issues: concurrency, algorithmic defects...
- *Testing*: dynamic, confronts the software to test suites
 - Can catch dynamic defects
 - 30% to 50% of software cost devoted to testing
 - More time spent on validation than on construction!
 - Exhaustive testing infeasible
 - **Testing can only show the presence of errors, not their absence!**

Hardware and software verification (3/4)

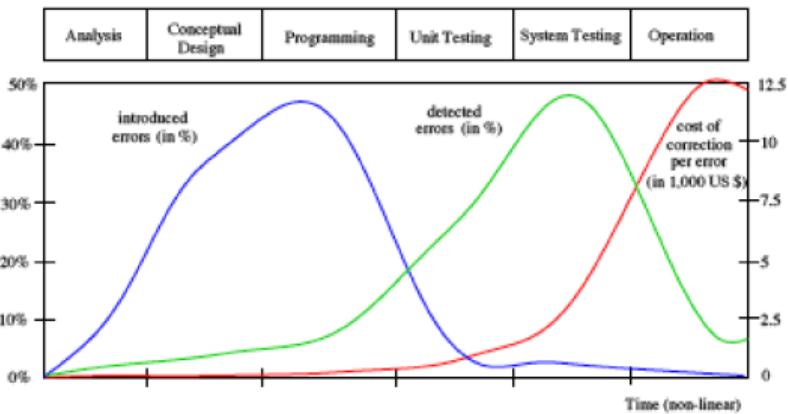
Catching bugs: the sooner, the better



Software lifecycle: error introduction, detection and repair costs

Hardware and software verification (3/4)

Catching bugs: the sooner, the better



Software lifecycle: error introduction, detection and repair costs

⇒ We need methods that can detect bugs early in a software's life.

Hardware and software verification (4/4)

Hardware verification

- Preventing errors is vital:
 - high fabrication costs,
 - fixing defects after delivery is difficult (no patch),
 - high quality expectations.

Hardware and software verification (4/4)

Hardware verification

- Preventing errors is vital:
 - high fabrication costs,
 - fixing defects after delivery is difficult (no patch),
 - high quality expectations.
- >50% of Application-Specific Integrated Circuits do not work properly after initial design and fabrication
- >70% of the total development time is devoted to error detection and prevention

Hardware and software verification (4/4)

Hardware verification

- Preventing errors is vital:
 - high fabrication costs,
 - fixing defects after delivery is difficult (no patch),
 - high quality expectations.
- >50% of Application-Specific Integrated Circuits do not work properly after initial design and fabrication
- >70% of the total development time is devoted to error detection and prevention
- Some techniques: *emulation* (~ testing), *simulation* (~ testing executed on models), *hardware testing* (to find fabrication faults)

Formal verification

Goal

Given

- a *formal model* of the system (= how it behaves)
- and a *formal specification* (= what it should do **and not do**),

check that the system satisfies the specification by (semi-)automatically generating some sort of *mathematical proof*.

Formal verification

Goal

Given

- a *formal model* of the system (= how it behaves)
- and a *formal specification* (= what it should do **and not do**),

check that the system satisfies the specification by (semi-)automatically generating some sort of *mathematical proof*.

Usefulness

- Early integration of verification in the design process.
- More effective verification (higher coverage).
- Reduced verification time.

Formal verification

Goal

Given

- a *formal model* of the system (= how it behaves)
- and a *formal specification* (= what it should do **and not do**),

check that the system satisfies the specification by (semi-)automatically generating some sort of *mathematical proof*.

Usefulness

- Early integration of verification in the design process.
- More effective verification (higher coverage).
- Reduced verification time.

⇒ safety ↗ and costs ↘

Specification formalisms

Formal encoding of such properties requires appropriate **specification formalisms**

Specification formalisms

Formal encoding of such properties requires appropriate **specification formalisms**

- Most are **temporal logics** (LTL, CTL, etc).
- *Not all logics can express all properties!*

Specification formalisms

Formal encoding of such properties requires appropriate **specification formalisms**

- Most are **temporal logics** (LTL, CTL, etc).
- *Not all logics can express all properties!*

Trade-off between **expressiveness** and **tractability**.

→ think about *decidability* and *complexity*: e.g., no hope of checking termination for Turing-powerful models.

Limits of formal verification (cf. validation)

Is the model right?

- Is it a faithful representation of the implementation?

Limits of formal verification (cf. validation)

Is the model right?

- Is it a faithful representation of the implementation?

Is the specification right?

- Often difficult to formalize, from oral language to logical formulas
- Difficult to check: does it really represent the expected behaviour of the system?

Limits of formal verification (cf. validation)

Is the model right?

- Is it a faithful representation of the implementation?

Is the specification right?

- Often difficult to formalize, from oral language to logical formulas
- Difficult to check: does it really represent the expected behaviour of the system?

Is the specification complete?

- Are all important properties specified?

Three approaches to verification (1/2)

Deductive methods (logical inference)

- Method: provide a formal **proof** that the property holds
- Tools: theorem provers and proof assistants/checkers (e.g., HOL, Isabelle)
- Applicable if the system has the form of a *mathematical theory*

Three approaches to verification (1/2)

Deductive methods (logical inference)

- Method: provide a formal **proof** that the property holds
- Tools: theorem provers and proof assistants/checkers (e.g., HOL, Isabelle)
- Applicable if the system has the form of a *mathematical theory*

Model-based simulation/testing

- Method: test the property by *exploring possible behaviours* of the model
- Applicable if the system defines an *executable model*

Three approaches to verification (2/2)

Model checking

- Method: systematic check of the property in all states of the model
- Tools: model checkers (e.g., Spin, NuSMV, Uppaal)
- Applicable if the system generates a *finitely representable behavioural model*
- Efficient techniques and tools
- If the property is not satisfied, can provide *counter-examples* (thus guiding repairs)

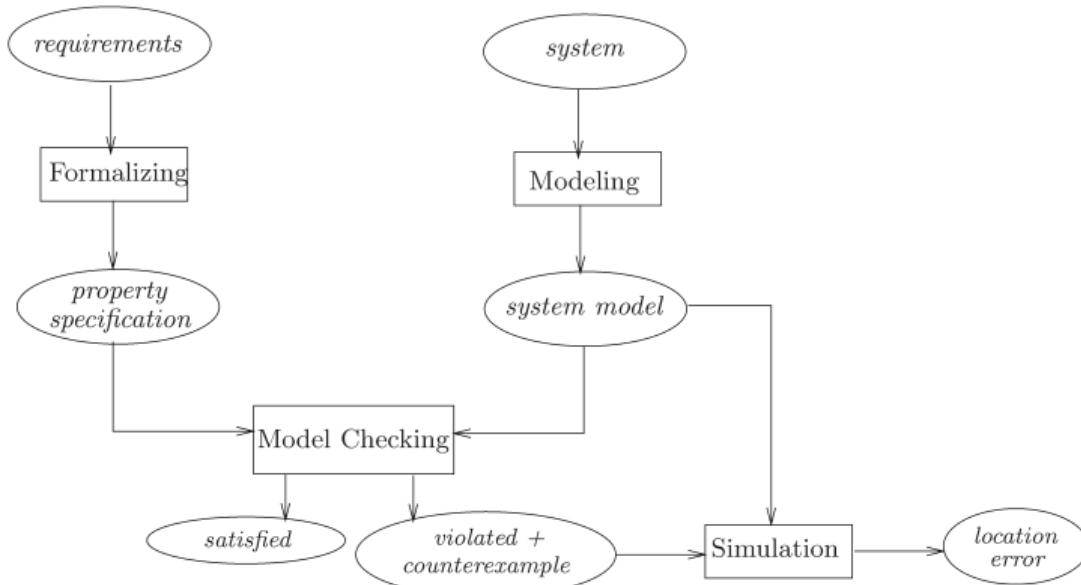
Three approaches to verification (2/2)

Model checking

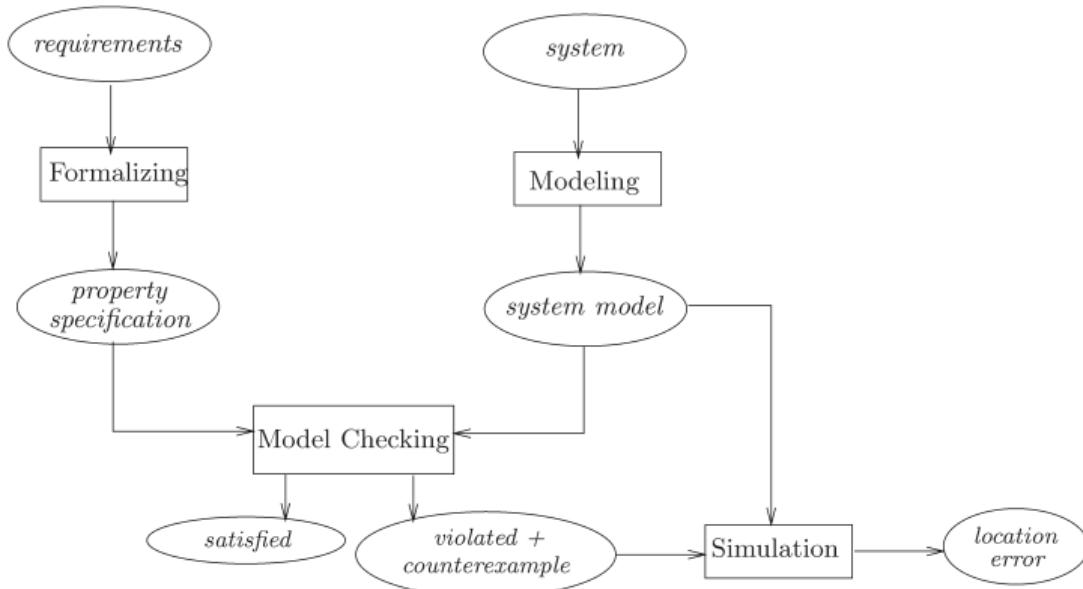
- Method: systematic check of the property in all states of the model
- Tools: model checkers (e.g., Spin, NuSMV, Uppaal)
- Applicable if the system generates a *finitely representable behavioural model*
- Efficient techniques and tools
- If the property is not satisfied, can provide *counter-examples* (thus guiding repairs)

↪ main focus of this course

Model checking process



Model checking process



Formal model
transition systems/Markov chains

Formal specification
LTL/CTL/PCTL

Pros of model checking

Pros:

- widely applicable (hardware, software, protocols),
- allows partial verification (most relevant properties),
- heavily automated,
- growing industrial interest,
- counter-example generation,
- sound mathematical foundations,
- not biased to the most probable scenarios (in contrast to *testing*)

Cons of model checking

Cons:

- focus on *control-intensive* applications (reactive systems)—less on *data-oriented* applications (batch processing systems),
- **model checking is only as good as the model,**
- decidability and complexity issues (state explosion problem),
- completeness is not guaranteed (if the specification omits important properties)

Cons of model checking

Cons:

- focus on *control-intensive* applications (reactive systems)—less on *data-oriented* applications (batch processing systems),
- **model checking is only as good as the model,**
- decidability and complexity issues (state explosion problem),
- completeness is not guaranteed (if the specification omits important properties)

All in all:

a quite effective technique to expose design errors.

↪ interesting addition to most design processes

Industry usage

Model checking techniques are increasingly present in industrial design processes

Industry usage

Model checking techniques are increasingly present in industrial design processes

- *Security*. A flaw in the Needham-Schroeder public-key protocol remained undiscovered for 17 years before being revealed by model checking

Industry usage

Model checking techniques are increasingly present in industrial design processes

- *Security*. A flaw in the Needham-Schroeder public-key protocol remained undiscovered for 17 years before being revealed by model checking
- *Model checkers for C, C++ and Java*; Developed and used by Microsoft, Digital, NASA; Successfully applied to the design of *device drivers*

Industry usage

Model checking techniques are increasingly present in industrial design processes

- *Security*. A flaw in the Needham-Schroeder public-key protocol remained undiscovered for 17 years before being revealed by model checking
- *Model checkers for C, C++ and Java*; Developed and used by Microsoft, Digital, NASA; Successfully applied to the design of *device drivers*
- **Facebook** has the verification tool Infer and its own verification team
- **Amazon** does verification within their automated-reasoning team

Awards for model checking advances



E. Clarke

A. Emerson

J. Sifakis

M. Vardi

P. Wolper

— Turing Award 2007 —

Gödel Prize 2000

Awards for model checking advances



E. Clarke

A. Emerson

J. Sifakis

M. Vardi

P. Wolper

— Turing Award 2007 —

Gödel Prize 2000

- Turing award: “for their role in developing model checking into a highly effective verification technology, widely adopted in the hardware and software industries.”

Outline

1 Course particularities

2 Why should we worry?

3 Formal verification in a nutshell

4 Conclusion & outline

Outline of the course

1. Introduction
2. Transition systems
3. Linear temporal logic
4. Omega automata
5. Binary decision diagrams - Proj. 1
6. Model checking LTL
7. Reactive synthesis
8. Parity games - Proj. 2